



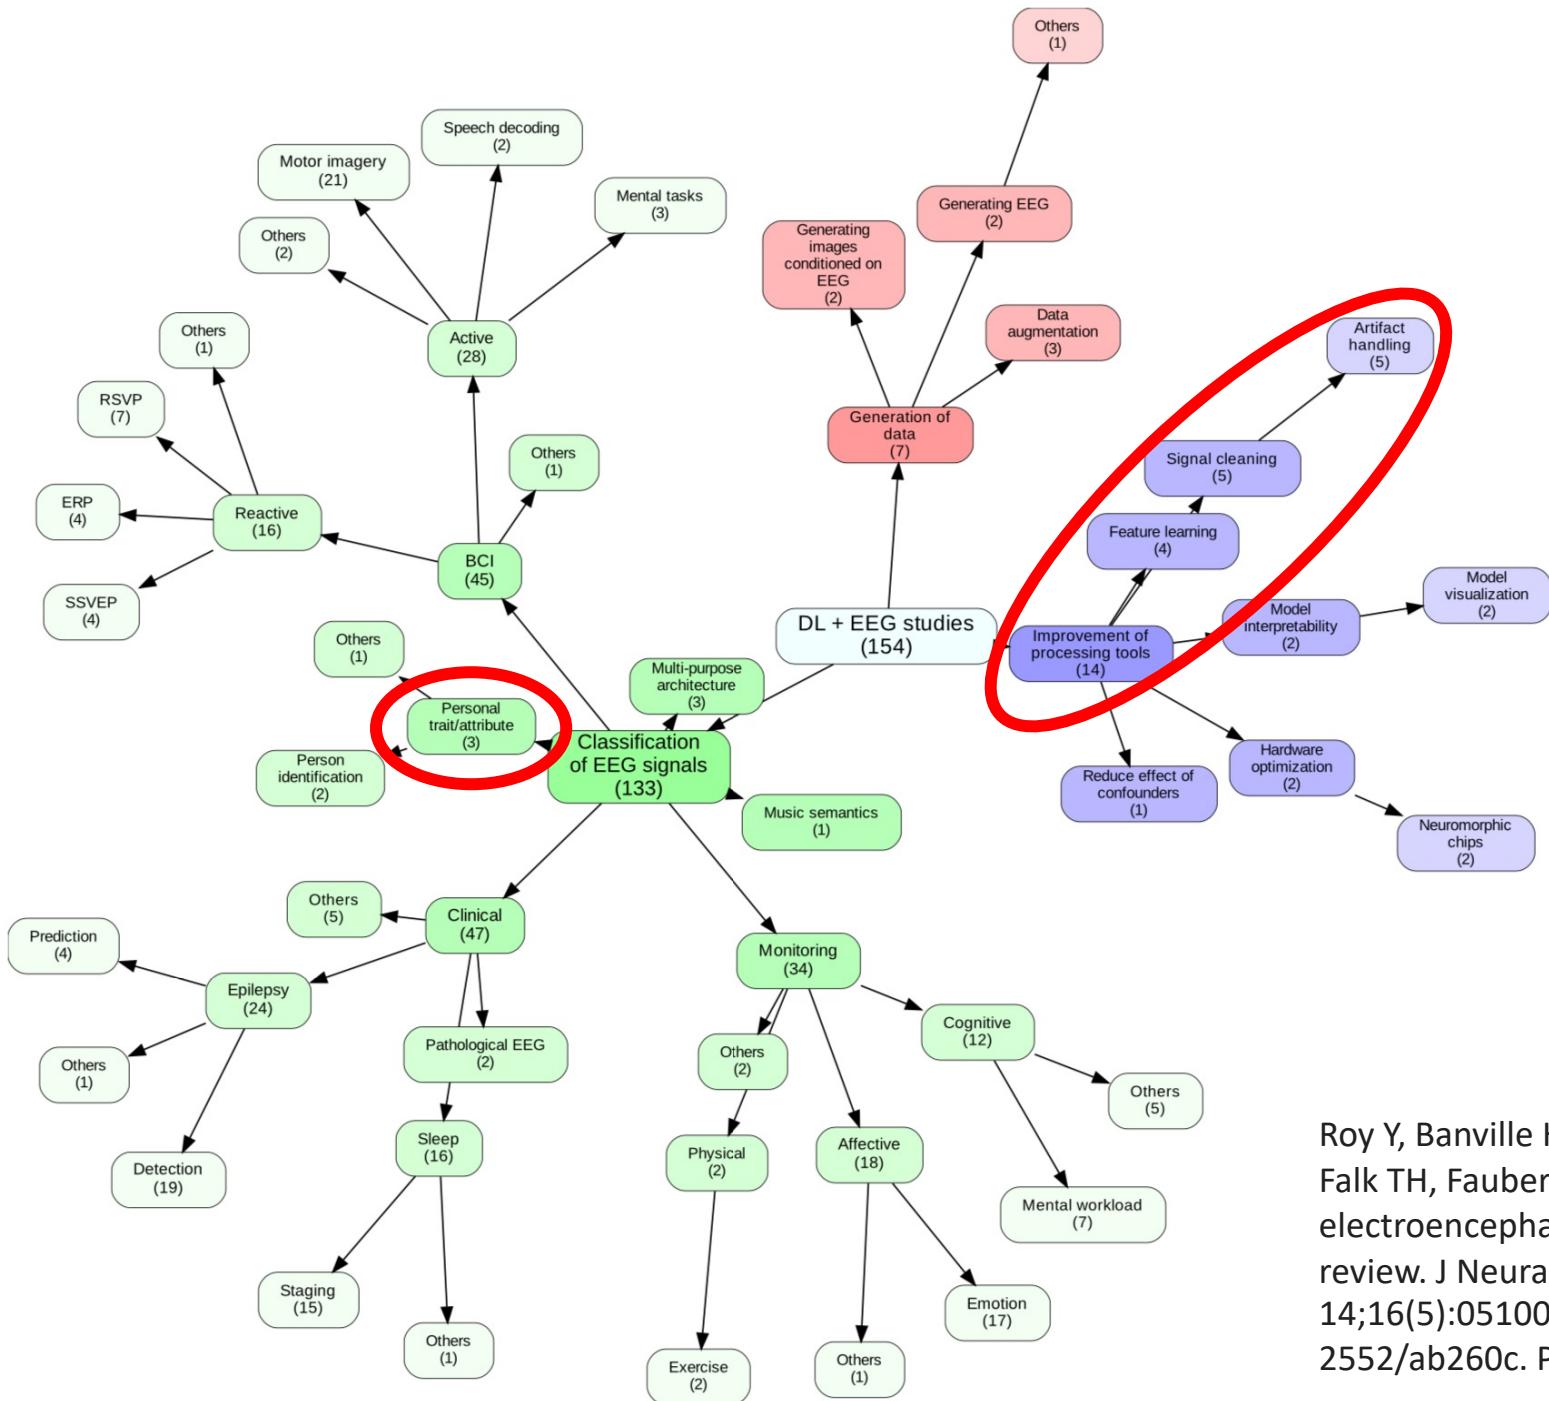
CerCo
UMR5549



Deep Learning applied to EEG

Arnaud Delorme & Dung Truong





Roy Y, Banville H, Albuquerque I, Gramfort A, Falk TH, Faubert J. Deep learning-based electroencephalography analysis: a systematic review. *J Neural Eng.* 2019 Aug 14;16(5):051001. doi: 10.1088/1741-2552/ab260c. PMID: 31151119.

Data Data Data Data Data Data Data Data



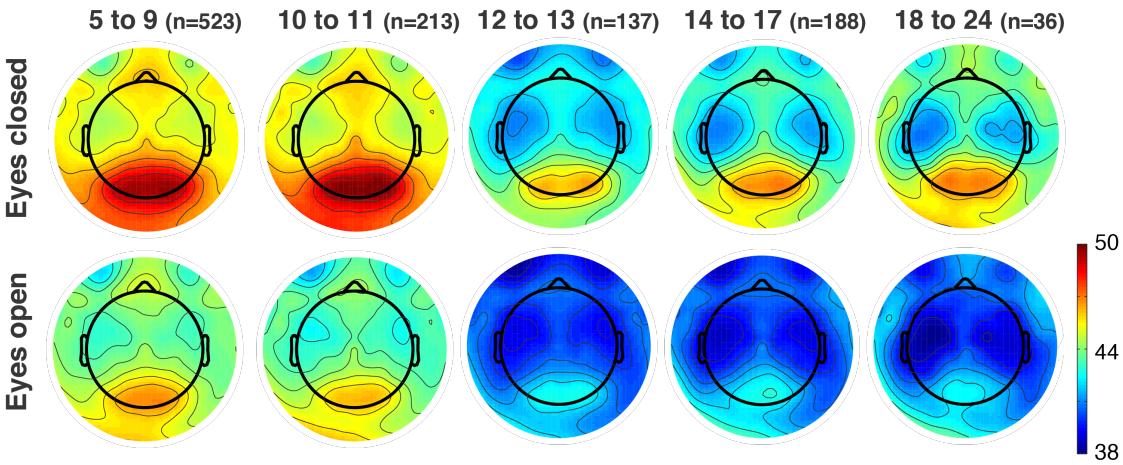
Michael Milham, PI

- 128 channels EEG
- ~3,000 EEG datasets (planed 10,000)
- Tasks involving emotions (The Gift movie)
- Rest (eyes open and eyes closed)



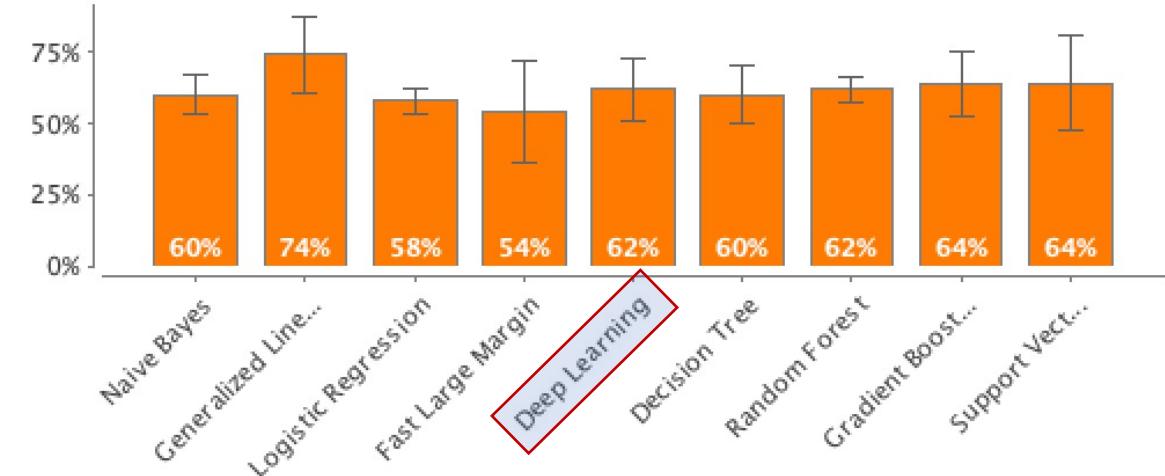
Eyes open vs closed

Different age groups eyes' open and eyes' closed from Child Mind Database (n=1,097)



Rapid Miner (3 hidden layers 100, 200, 100)

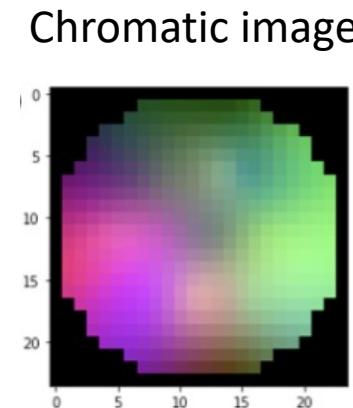
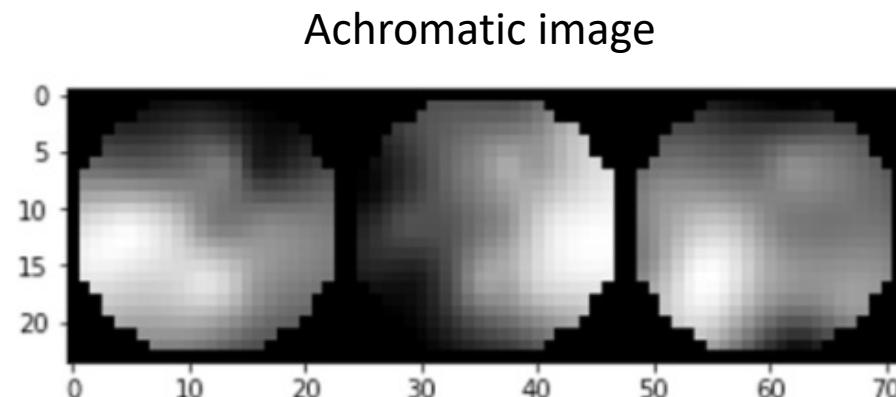
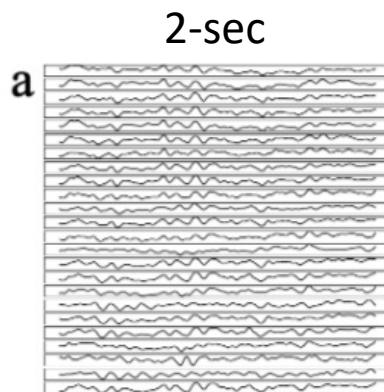
Accuracy



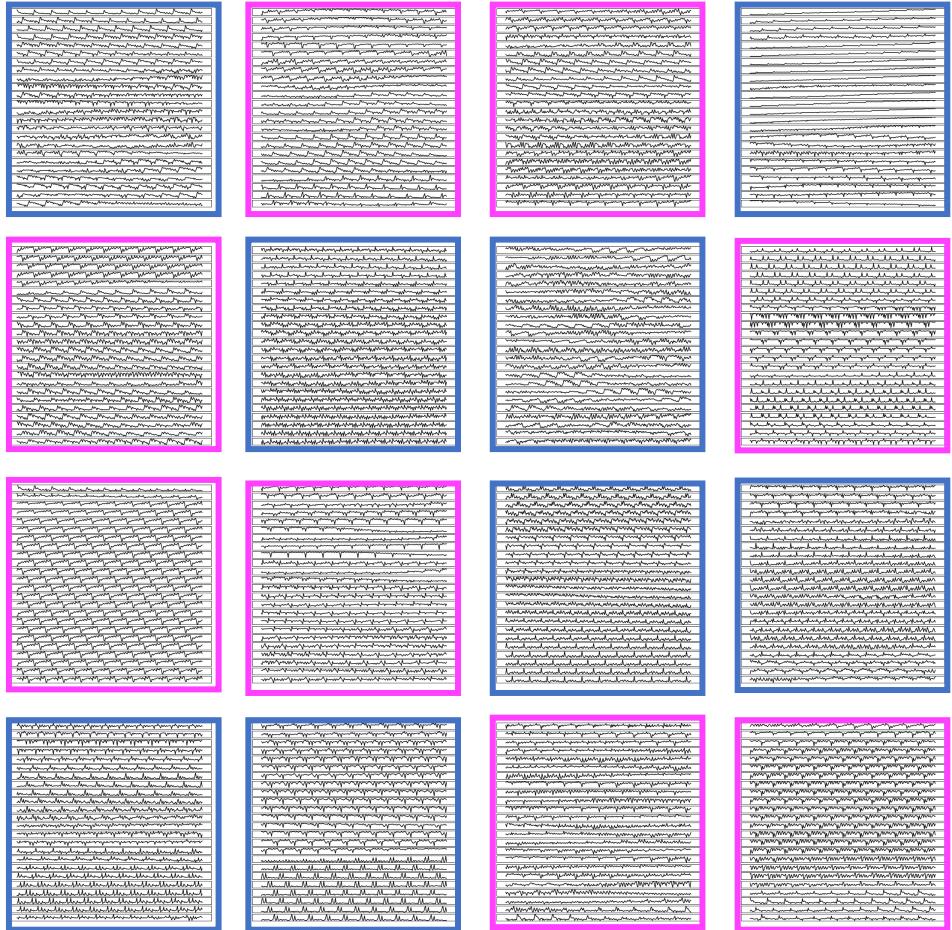
Machine learning/Deep learning

- 1574 participants (50% female) 128-channels
- 2-second extracted epochs eyes open and eyes closed
- 5 predictors: **sex**, **handedness**, **eyes open/closed**, **age**, **segment count**

↑
Categorical
↑
Continuous



Input data 24 x 256 x n



Training

944 subjects for training set (50% male)

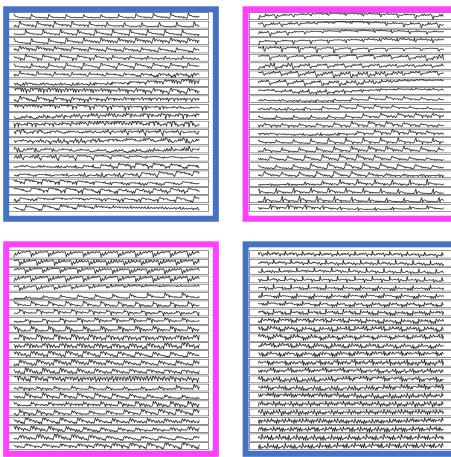
472 subjects for validation set (50% male)

Testing

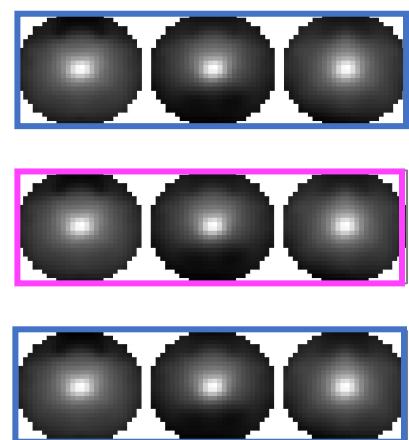
198 subjects for set 1 (50% male)

650 subjects for set 2 (100% male)

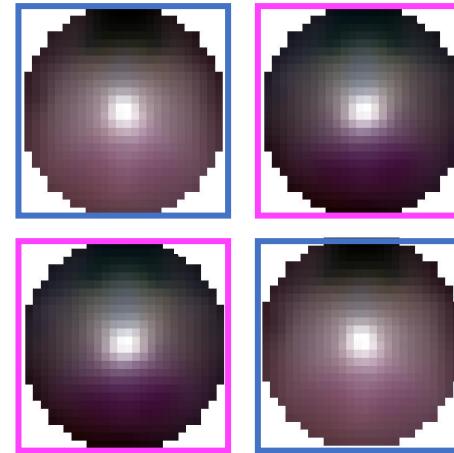
Raw input data
24 x 256 x n



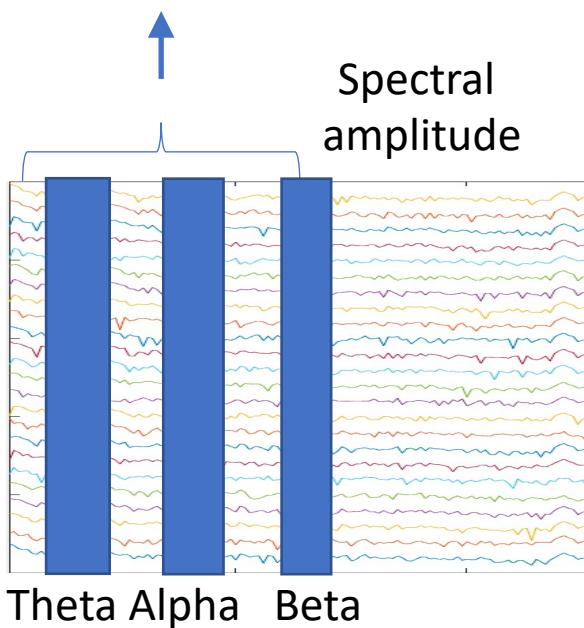
Spectral data
24 x 72 x n

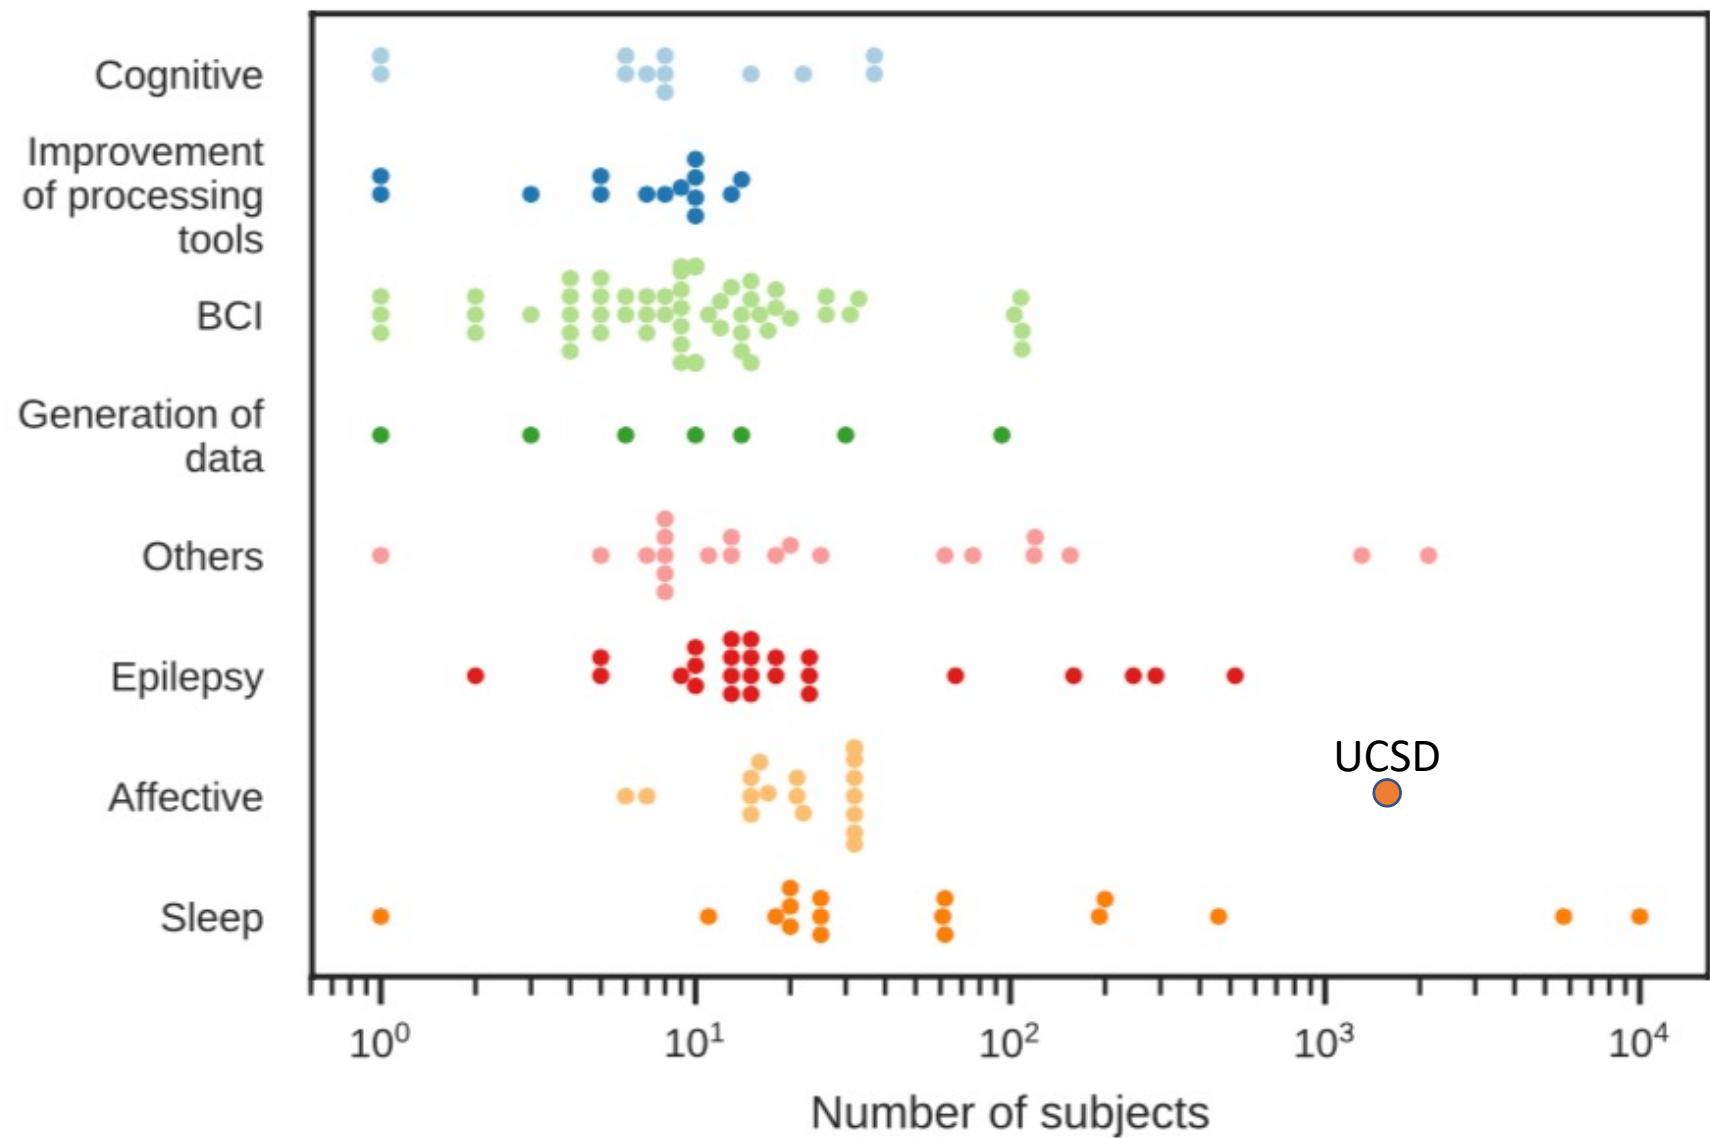


Spectral data
24 x 24 x n



theta alpha beta





Sex classification

www.nature.com/scientificreports/

SCIENTIFIC REPORTS



OPEN

Predicting sex from brain rhythms with deep learning

Michel J. A. M. van Putten¹, Sebastian Olbrich² & Martijn Arns³

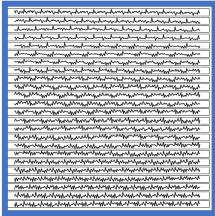
Received: 12 October 2017
Accepted: 6 February 2018
Published online: 15 February 2018

We have excellent skills to extract sex from visual assessment of human faces, but assessing sex from human brain rhythms seems impossible. Using deep convolutional neural networks, with unique potential to find subtle differences in apparent similar patterns, we explore if brain rhythms from either sex contain sex specific information. Here we show, in a ground truth scenario, that a deep neural net can predict sex from scalp electroencephalograms with an accuracy of $>80\%$ ($p < 10^{-5}$), revealing that brain rhythms are sex specific. Further, we extracted sex-specific features from the deep net filter layers, showing that fast beta activity (20–25 Hz) and its spatial distribution is a main distinctive attribute. This demonstrates the ability of deep nets to detect features in spatiotemporal data unnoticed by visual assessment, and to assist in knowledge discovery. We anticipate that this approach may also be successfully applied to other specialties where spatiotemporal data is abundant, including neurology, cardiology and neuropsychology.

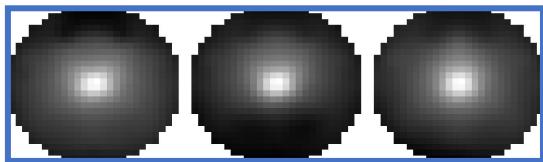
Dung Truong
MS Deep Learning



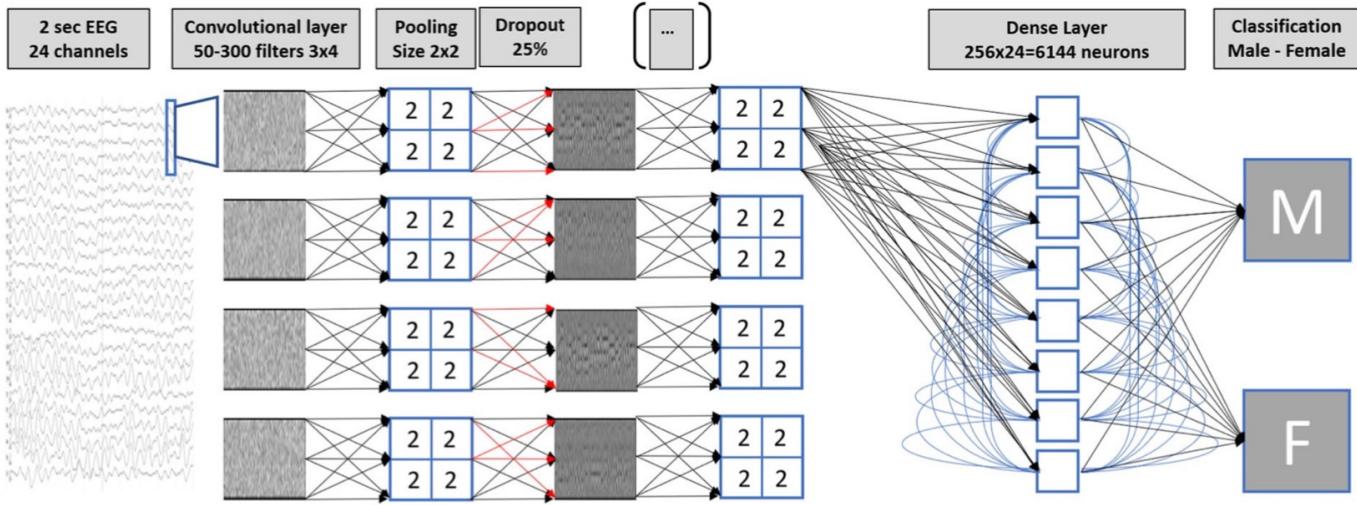
R-SCNN



S-SCNN



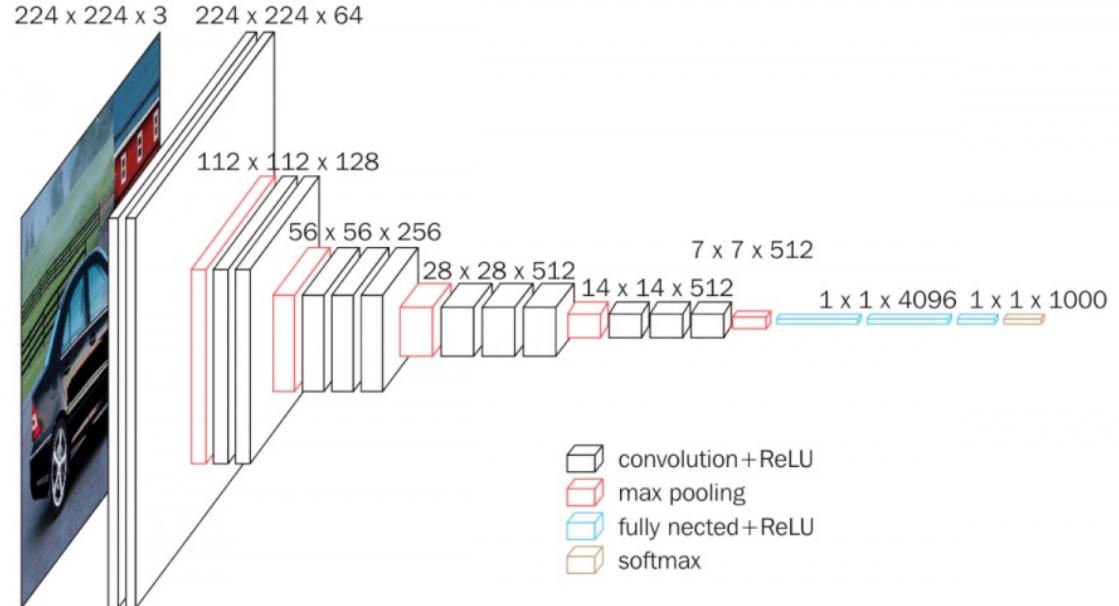
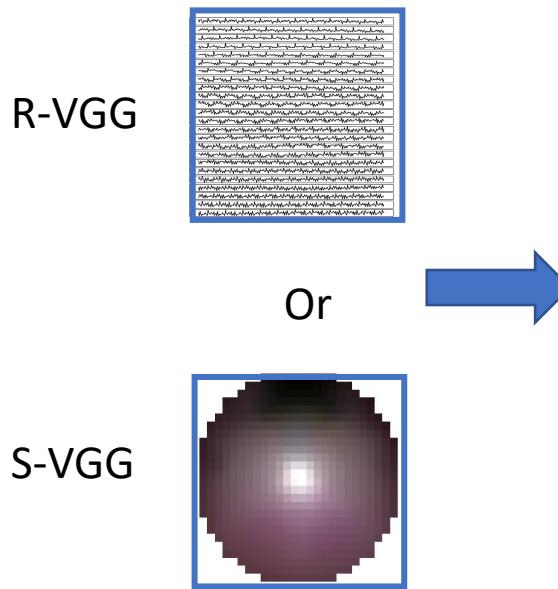
Or



Layer	Filter size	# of filters/hidden units
Convolutional	3x3	100
MaxPooling		
Dropout (25%)		
Convolutional	3x3	100
MaxPooling		
Dropout (25%)		
Convolutional	2x3	300
MaxPooling		
Dropout (25%)		
Convolutional†	1x7	300
MaxPooling*		
Dropout (25%)		
Convolutional†	1x3	100
Convolutional†	1x3	100
Fully connected		6144
Fully connected		2
Softmax		

Modified VGG16

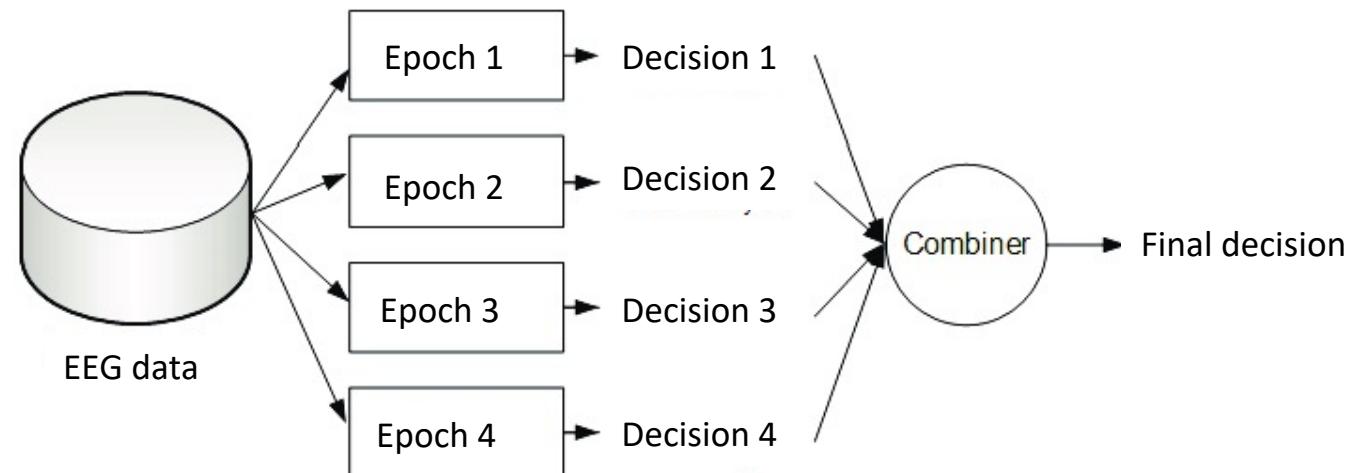
- Changed input size
- Keep the same scaling between layers
- Number of convolutions divided by 8 for each layer
- Dropped layers 19 to 32
- Change number of output classes to 2
- Retrain the whole network

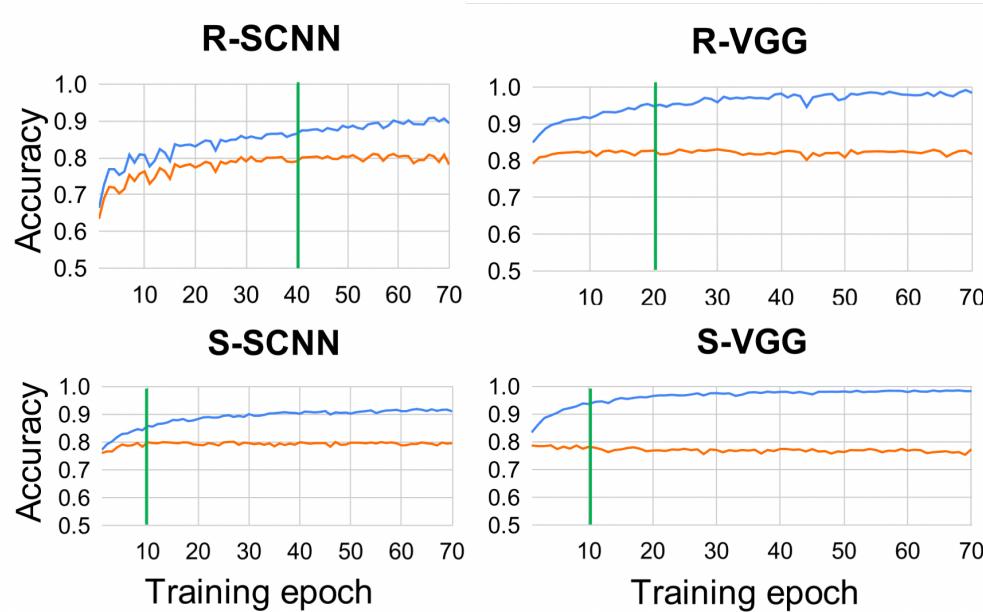


Majority voting



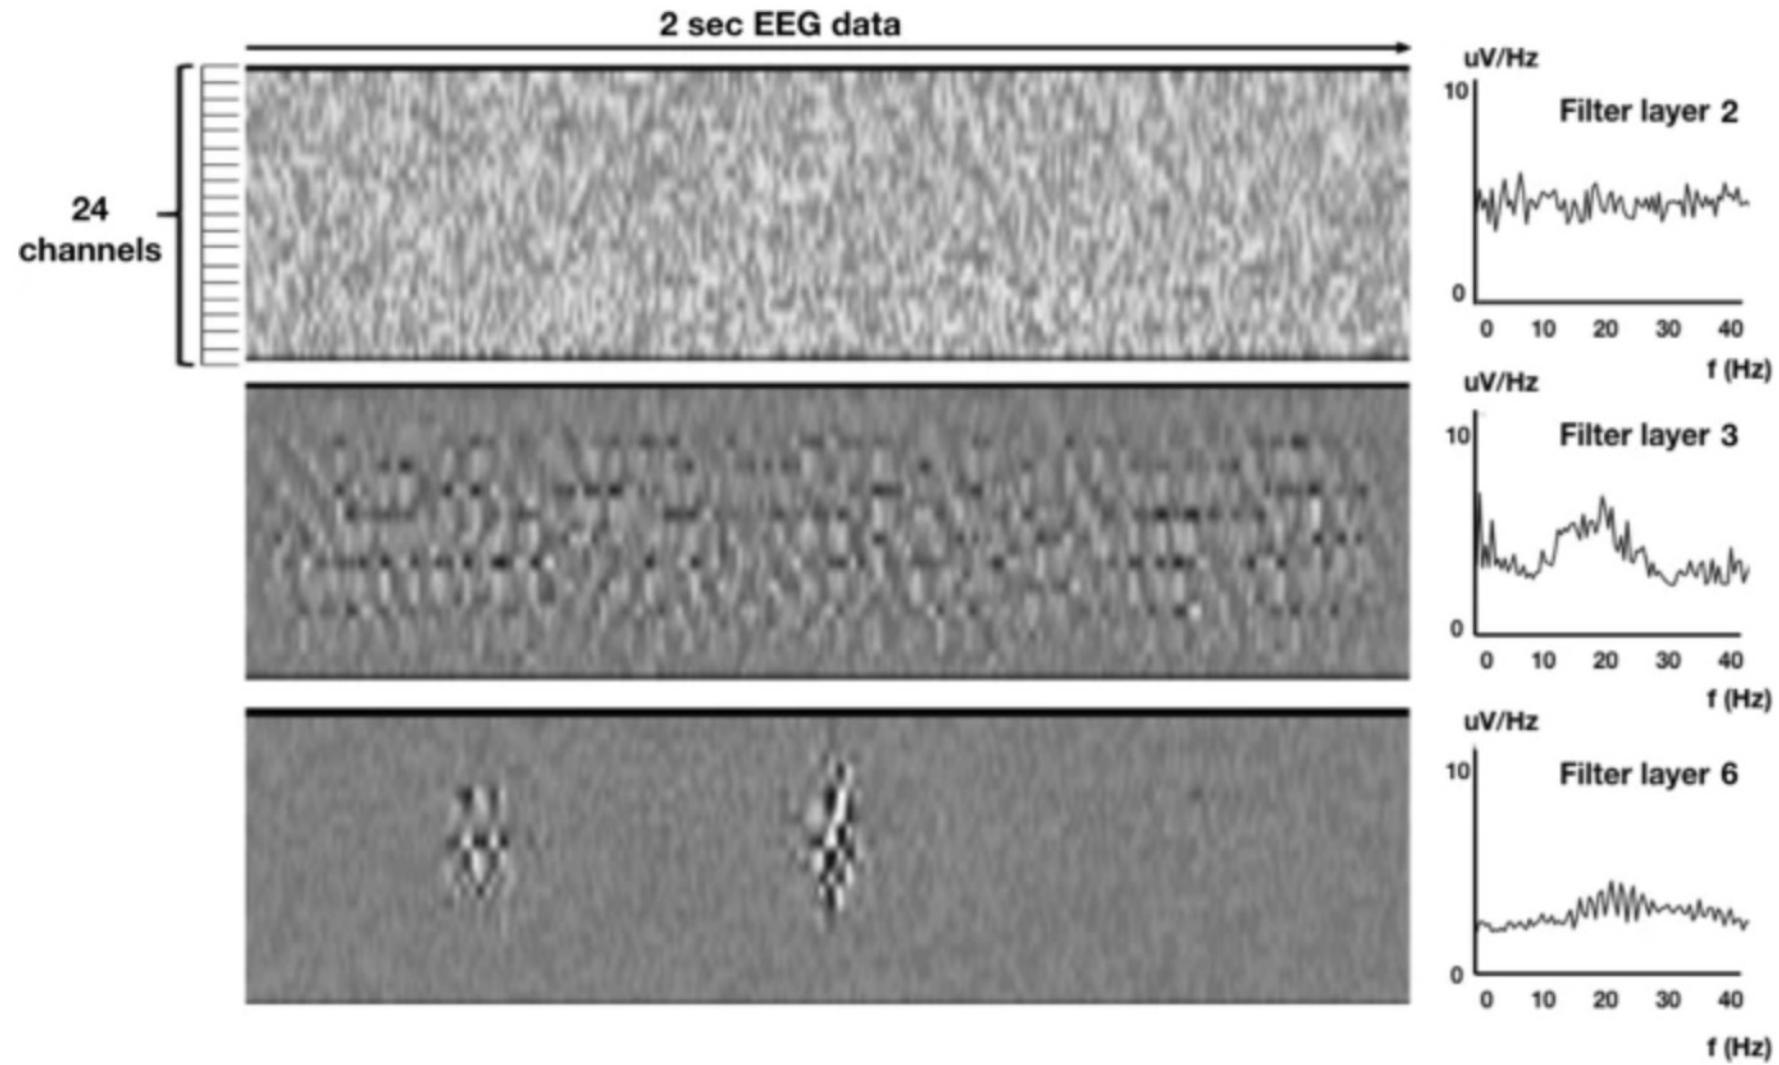
- Estimates are on 2-second epochs
- We have about 85 per subject
- Boost performance by about 8% (see next slide)





Model	Per-sample	Per-subject
	2-sec	2-sec x n with vote
R-SCNN	80.6 (79.7 to 81.5)	85.1 (84.3 to 85.9)
R-VGG	83.1 (82.7 to 83.4)	87.0 (86.6 to 87.4)
S-SCNN	79.0 (78.7 to 79.3)	83.2 (82.1 to 84.3)
S-VGG	77.1 (76.8 to 77.4)	81.3 (80.0 to 82.6)

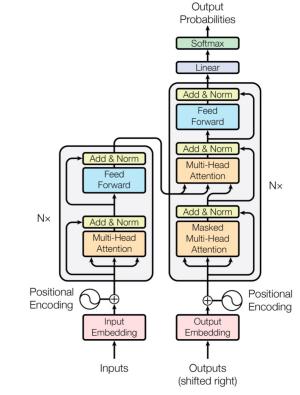
Table 3. Models classification accuracy. 95% confidence interval is indicated in parenthesis. Bolded values indicate best performance.



van Putten MJAM, Olbrich S, Arns M. Predicting sex from brain rhythms with deep learning.
Sci Rep. 2018 Feb 15;8(1):3069. doi: 10.1038/s41598-018-21495-7. PMID: 29449649; PMCID:
PMC5814426.

Future developments

- New interpretations/features for EEG
- New architectures
- Benchmark data and scripts (Matlab and Python)
- EEGLAB data exporter



Interpretation

Reconstruction

3/2/add_5

Type: Add
Channels: 1,280
Convolution: [1,1]

Technique

- Feature Visualization
- DeepDream
- Dataset Samples
- Caricature
- Text Feature Visualization

An artificial, optimized image that maximizes activations of the given unit. [Read more.](#)

Params

Optimization Objective

- channel
- neuron

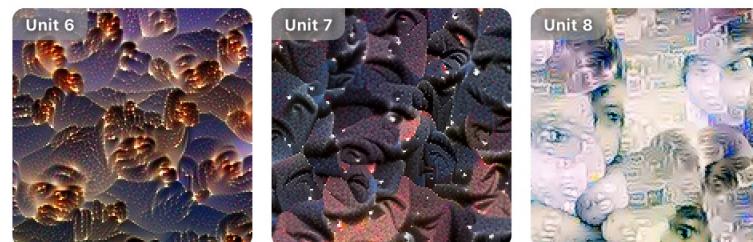
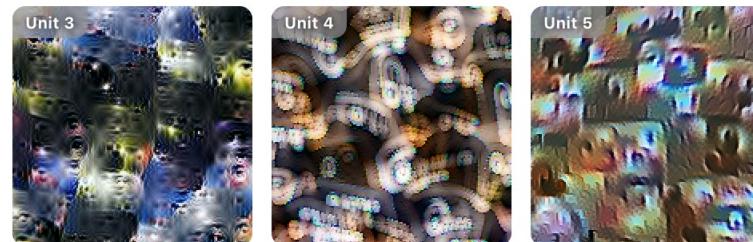
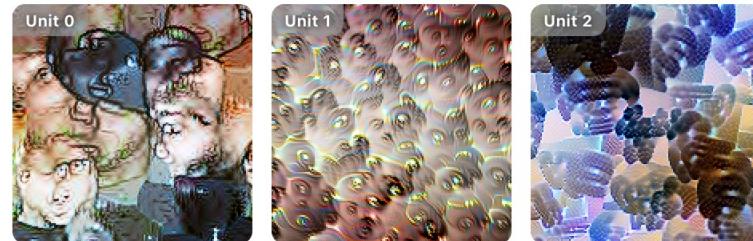
View

Image Size



Resize Behavior

- Crop image
- Scale image



Perspective

- Attention networks

Vaswani et al. (2017) All you need is attention. 31st Conference on Neural Information Processing Systems (NIPS 2017), Long Beach, CA, USA.

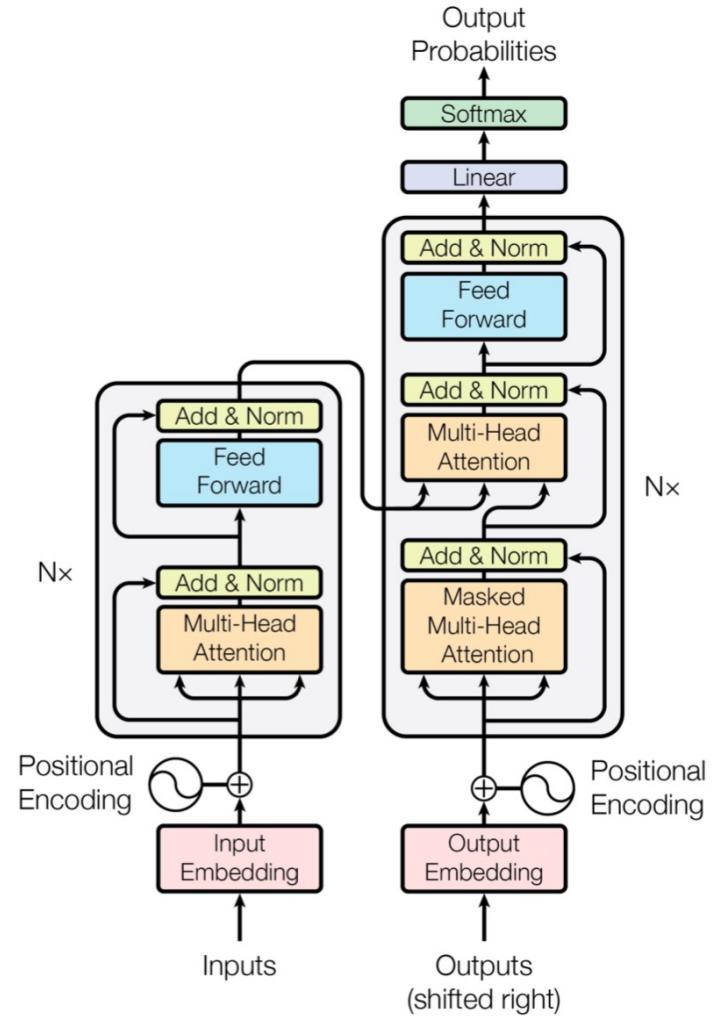
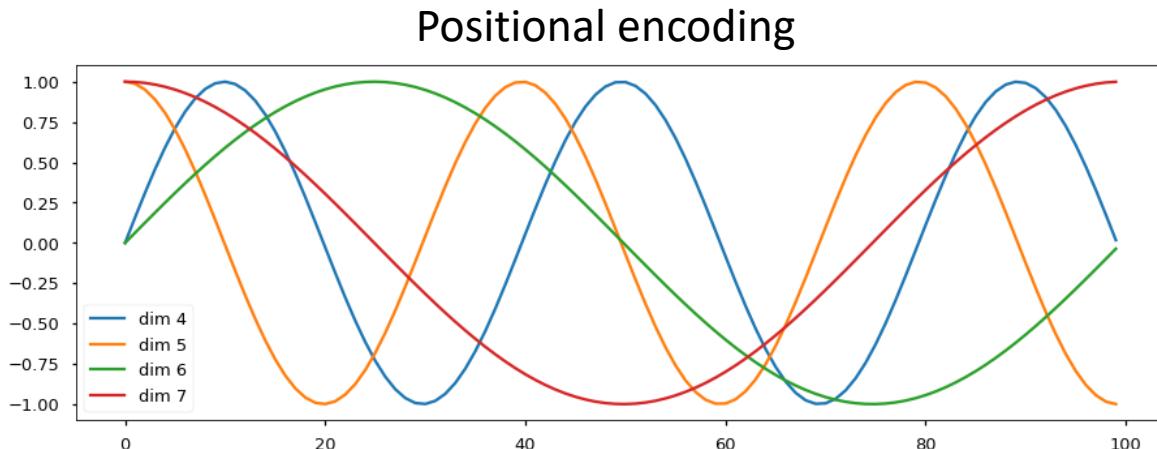


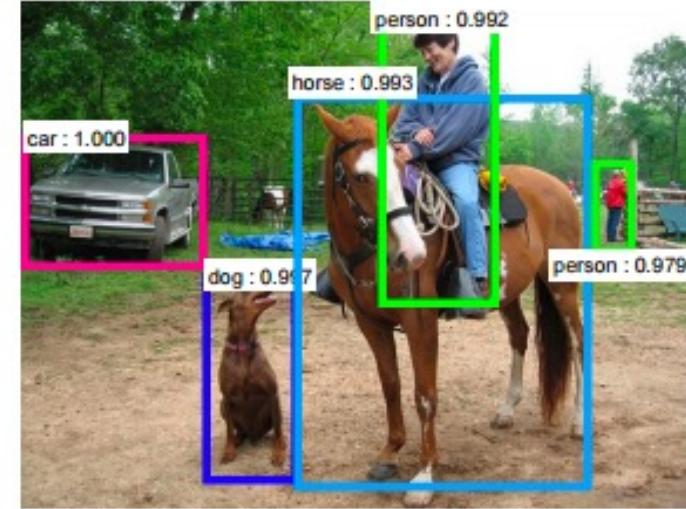
Figure 1: The Transformer - model architecture.



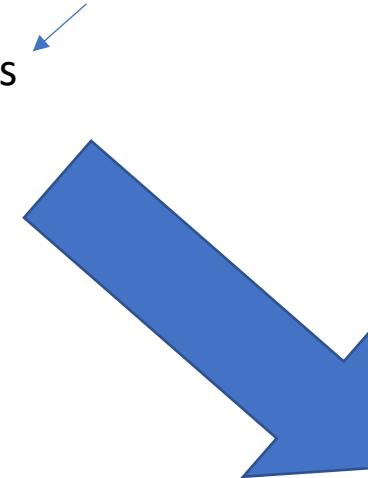
- EEG
- Other bio sensors
- Motion
- Eye tracker



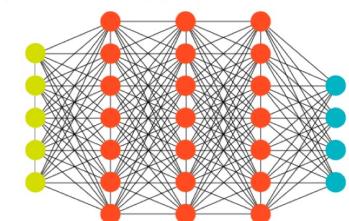
Event labels



- Automated object labeling
- Automated sound labeling

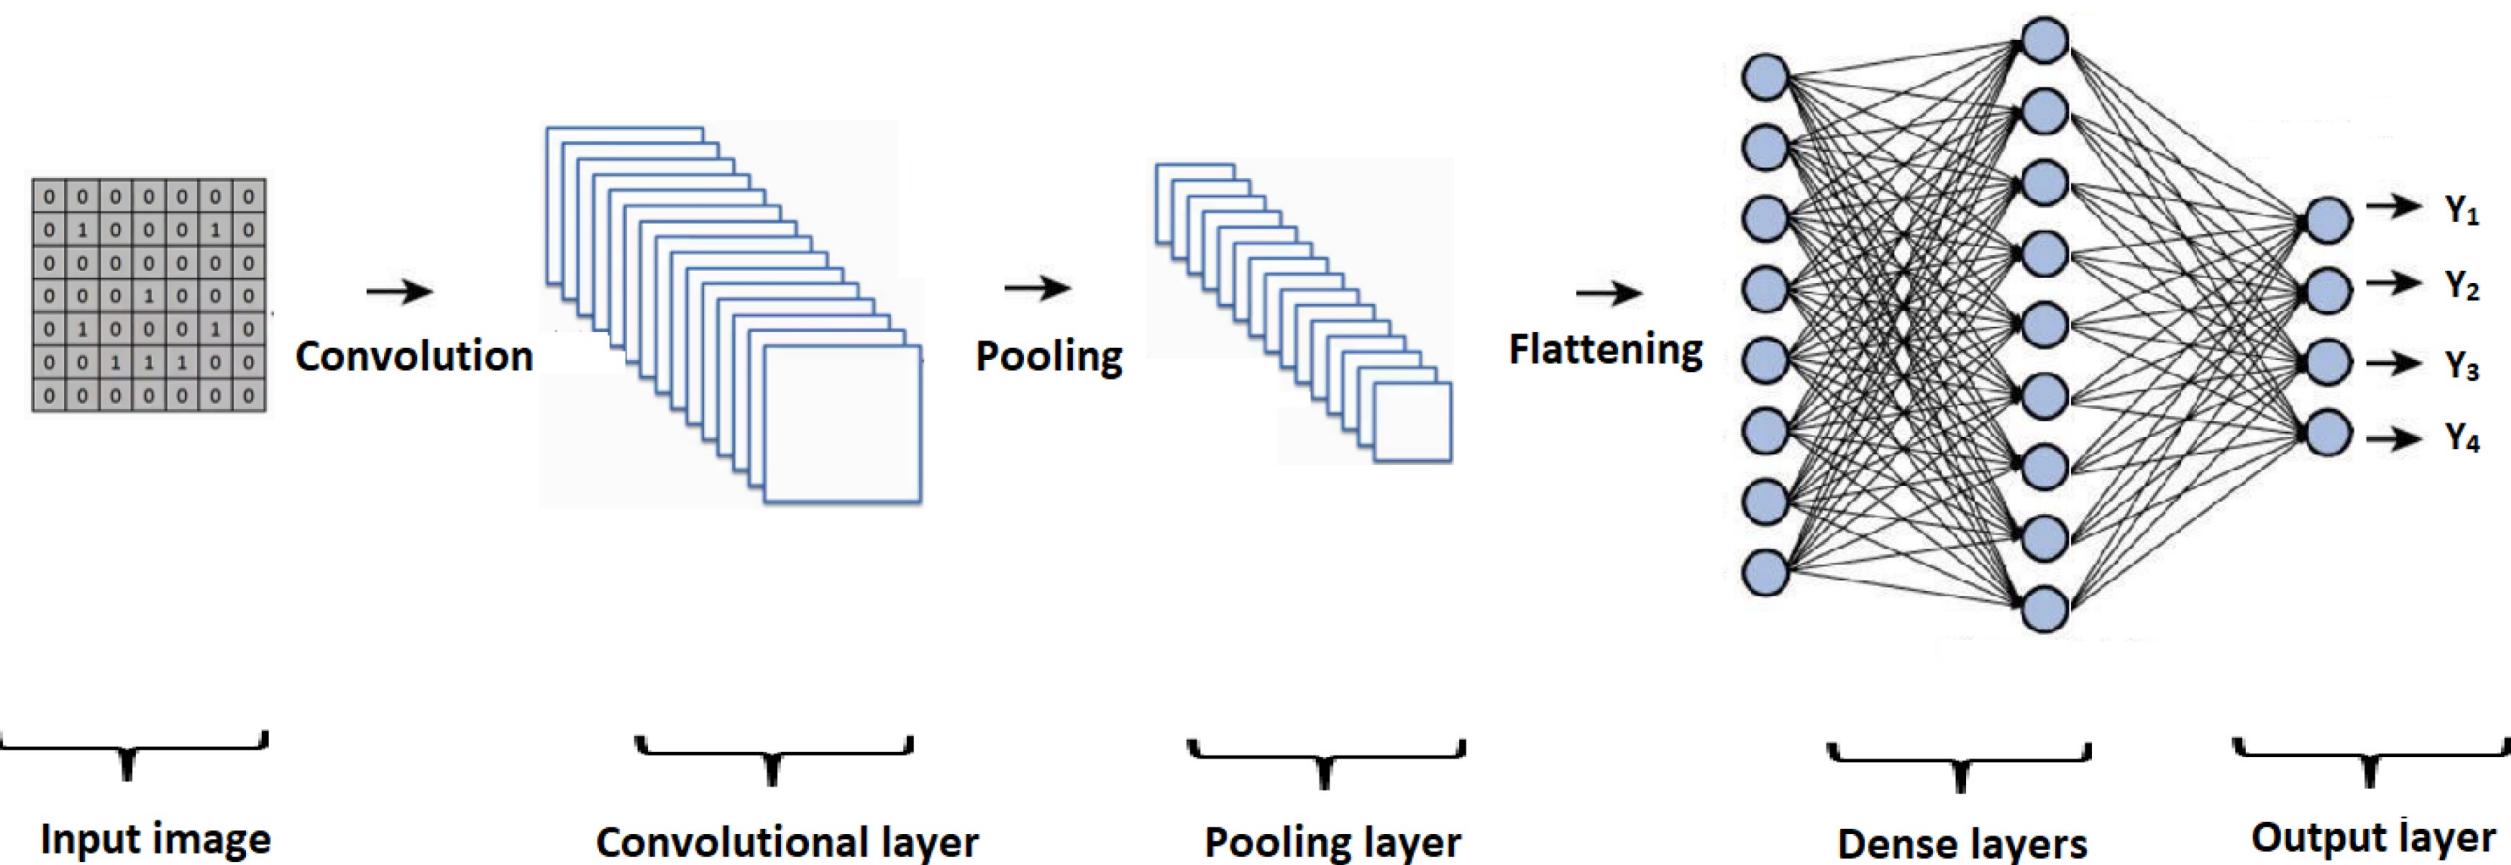


Deep learning

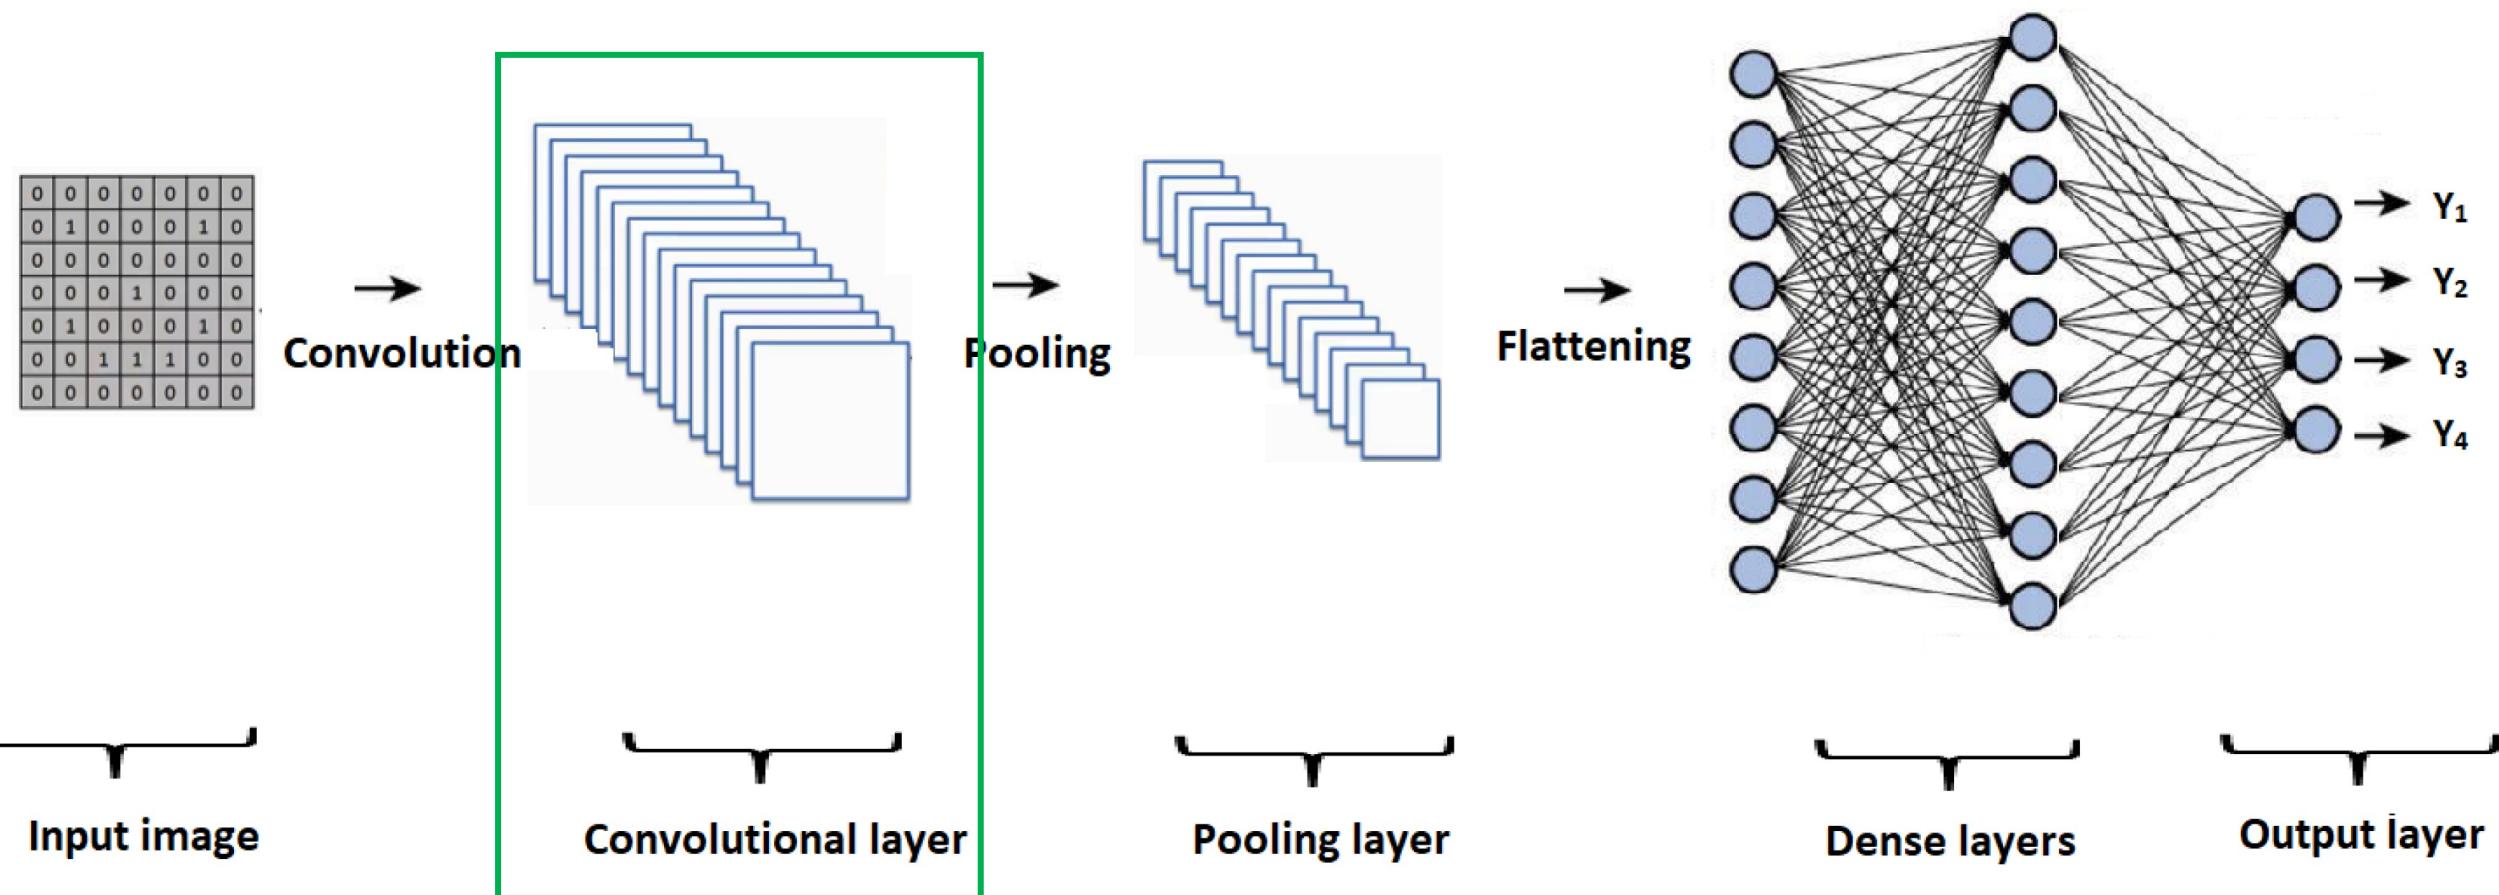




Convolutional Neural Network Architecture



Convolutional Neural Network Architecture



Convolutional layer

- Convolution operation

Input

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

*

Kernel/Filter

1	0	-1
1	0	-1
1	0	-1

=

Feature/Activation map

6		

$$7 \times 1 + 4 \times 1 + 3 \times 1 + \\ 2 \times 0 + 5 \times 0 + 3 \times 0 + \\ 3 \times -1 + 3 \times -1 + 2 \times -1 \\ = 6$$

Convolutional layer

- Convolution operation

Input

7	2	3	3	8
4	5	3	8	4
3	3	2	8	4
2	8	7	2	7
5	4	4	5	4

* Kernel/Filter

1	0	-1
1	0	-1
1	0	-1

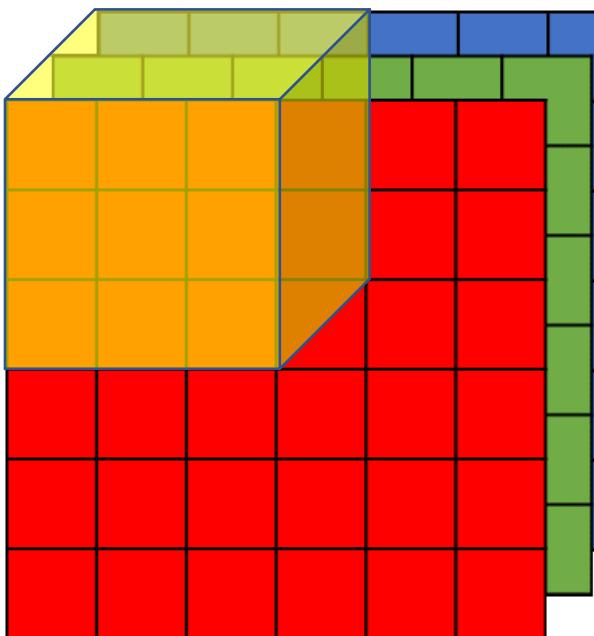
= Feature/Activation map

6		

$$7 \times 1 + 4 \times 1 + 3 \times 1 + \\ 2 \times 0 + 5 \times 0 + 3 \times 0 + \\ 3 \times -1 + 3 \times -1 + 2 \times -1 \\ = 6$$

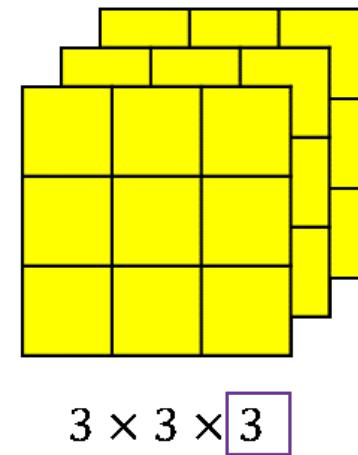
Convolutional layer

- Convolution over volume



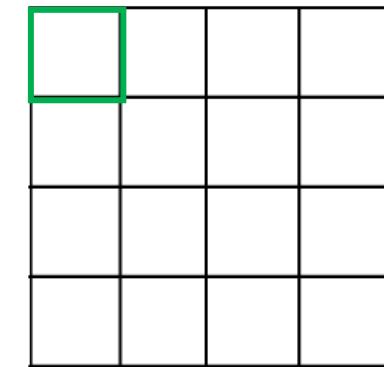
$6 \times 6 \times 3$

*



$3 \times 3 \times 3$

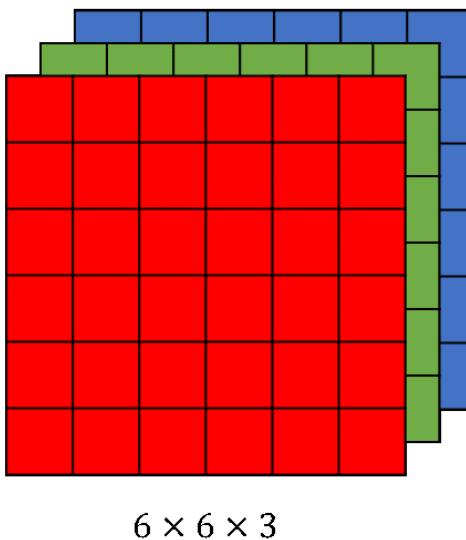
=



4×4

Convolutional layer

- Stacking filters



$$\begin{matrix} & \text{Vertical edge} \\ * & \begin{matrix} \text{3} \times \text{3} \times \text{3} \end{matrix} \\ & \text{Horizontal edge} \\ * & \begin{matrix} \text{3} \times \text{3} \times \text{3} \end{matrix} \end{matrix}$$

The diagram shows two convolution operations. The first operation takes the input volume and applies a "Vertical edge" filter, resulting in a feature map of size $3 \times 3 \times 3$. The second operation applies a "Horizontal edge" filter, also resulting in a feature map of size $3 \times 3 \times 3$.

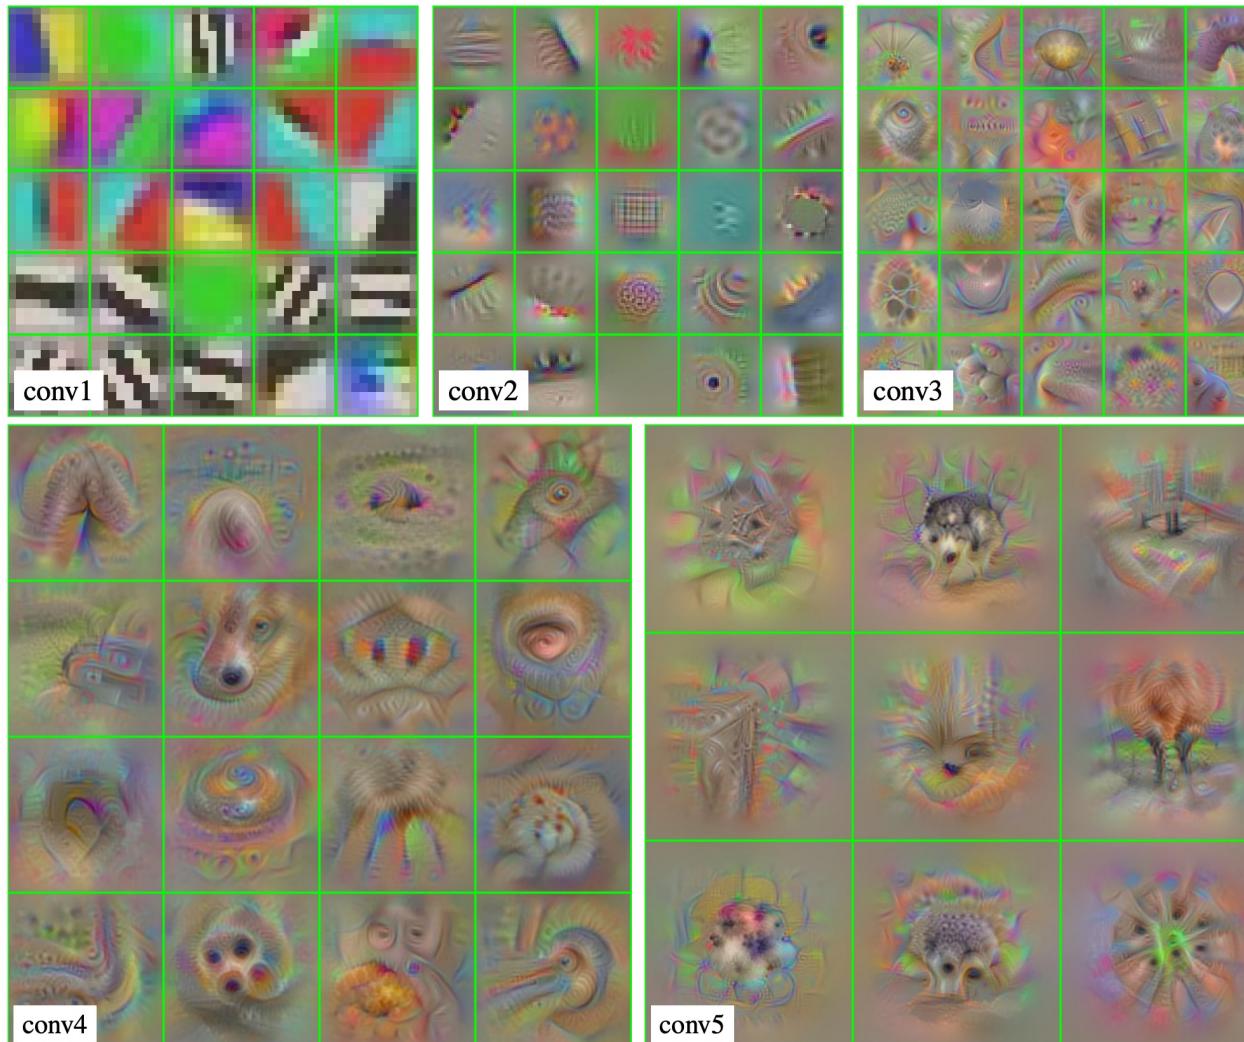
$$\begin{matrix} = & \begin{matrix} \text{4} \times \text{4} \times \text{2} \end{matrix} \\ = & \begin{matrix} \text{4} \times \text{4} \\ \text{4} \times \text{4} \times \text{2} \end{matrix} \end{matrix}$$

The diagram illustrates the result of stacking the two feature maps. The first feature map is a 4×4 grid of white squares. The second feature map is a $4 \times 4 \times 2$ volume represented as a cube. A curved arrow points from the second feature map to its equivalent representation as a 4×4 grid.

Convolutional layer



Convolutional layer



Convolutional layer

- Padding

0	0	0	0	0	0	0
0	60	113	56	139	85	0
0	73	121	54	84	128	0
0	131	99	70	129	127	0
0	80	57	115	69	134	0
0	104	126	123	95	130	0
0	0	0	0	0	0	0

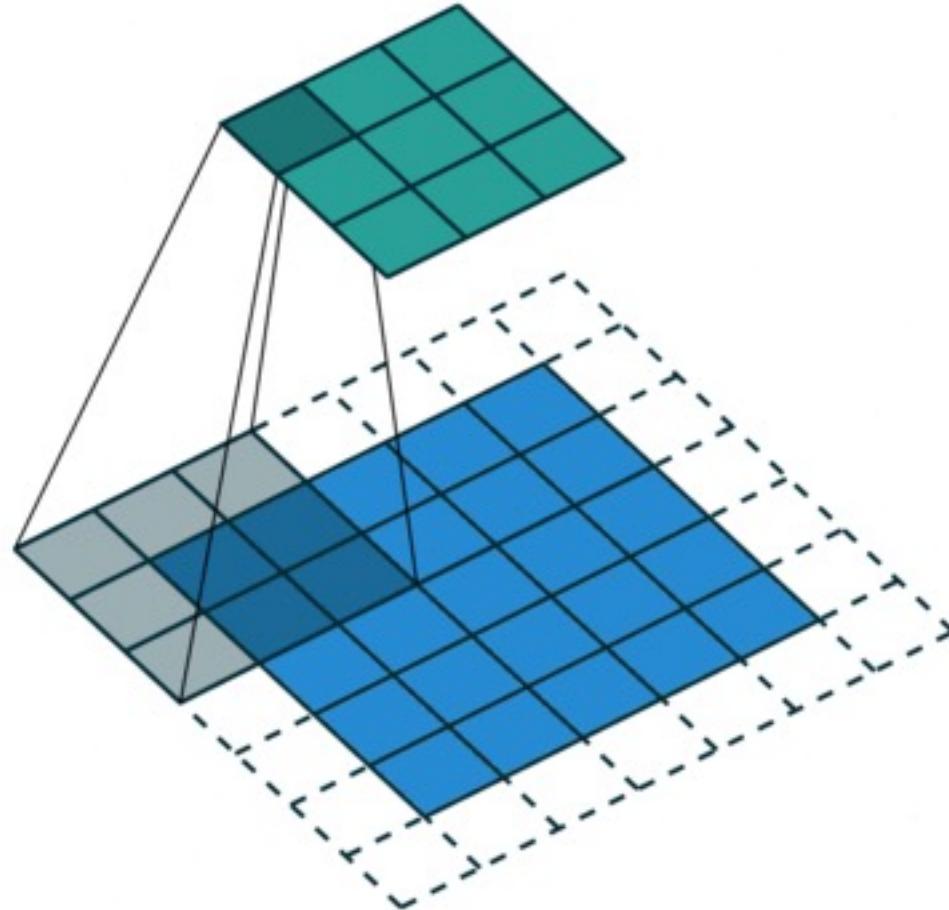
Kernel

0	-1	0
-1	5	-1
0	-1	0

114				

Convolutional layer

- Stride

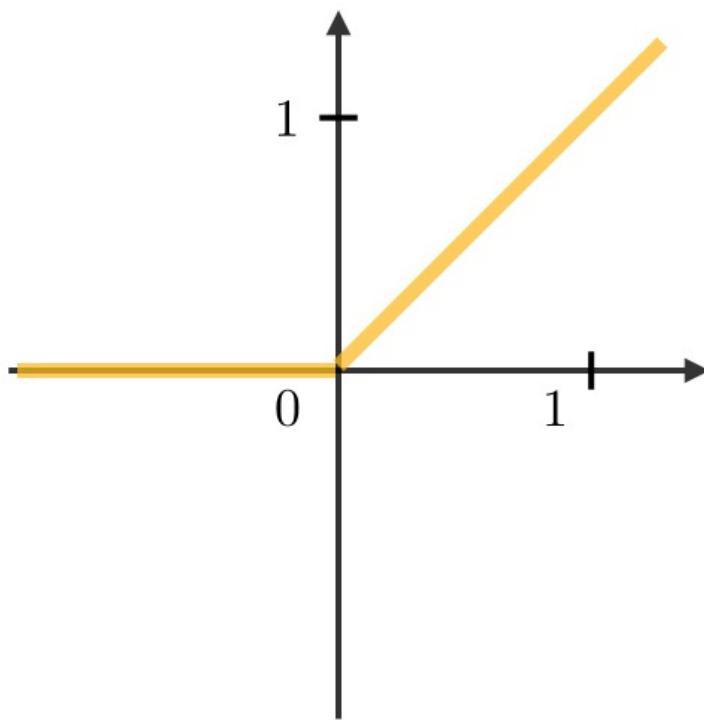


Example: Padding = 1, Stride = 2

Activation functions

- Rectified Linear Unit (ReLU)

$$g(z) = \max(0, z)$$



Filter 1 Feature Map

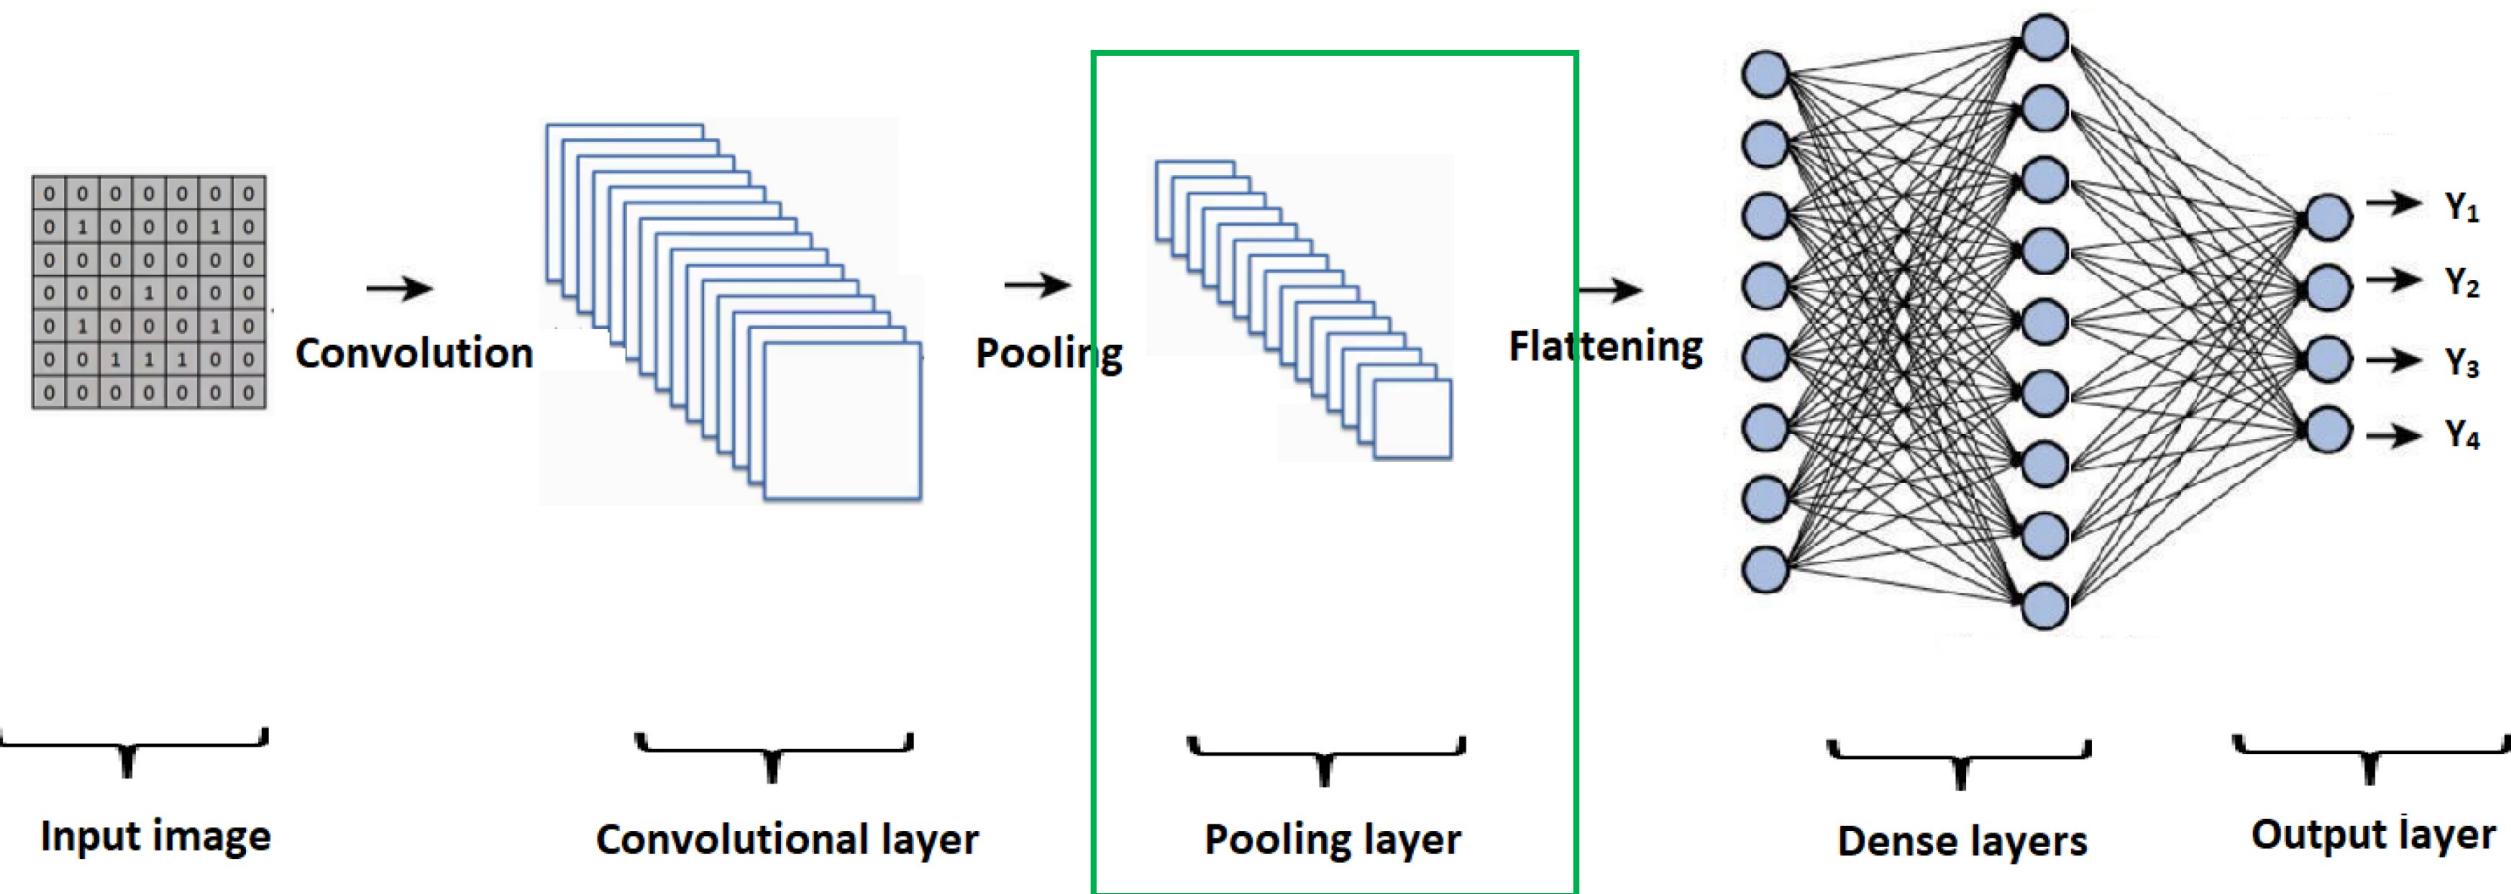
9	3	5	-8
-6	2	-3	1
1	3	4	1
3	-4	5	1

ReLU Layer

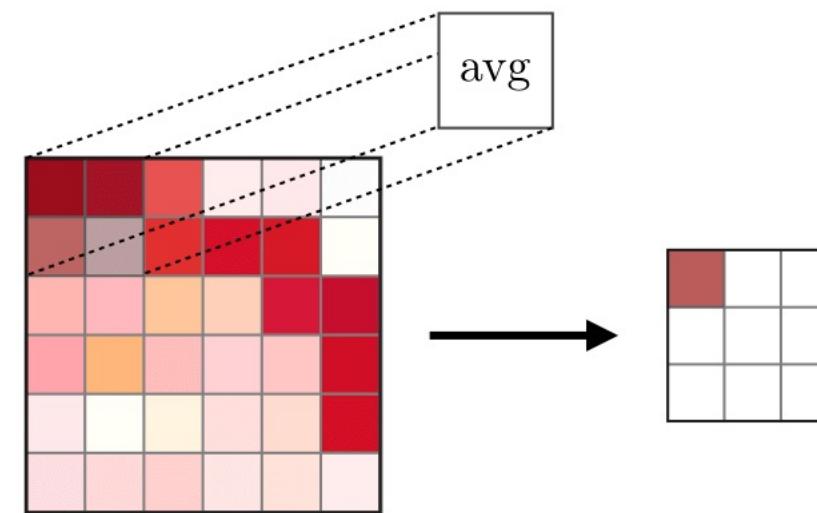
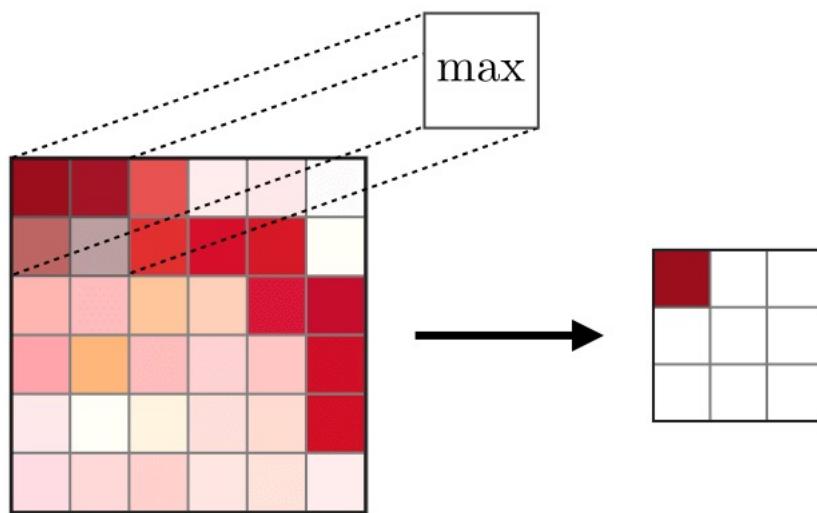


9	3	5	0
0	2	0	1
1	3	4	1
3	0	5	1

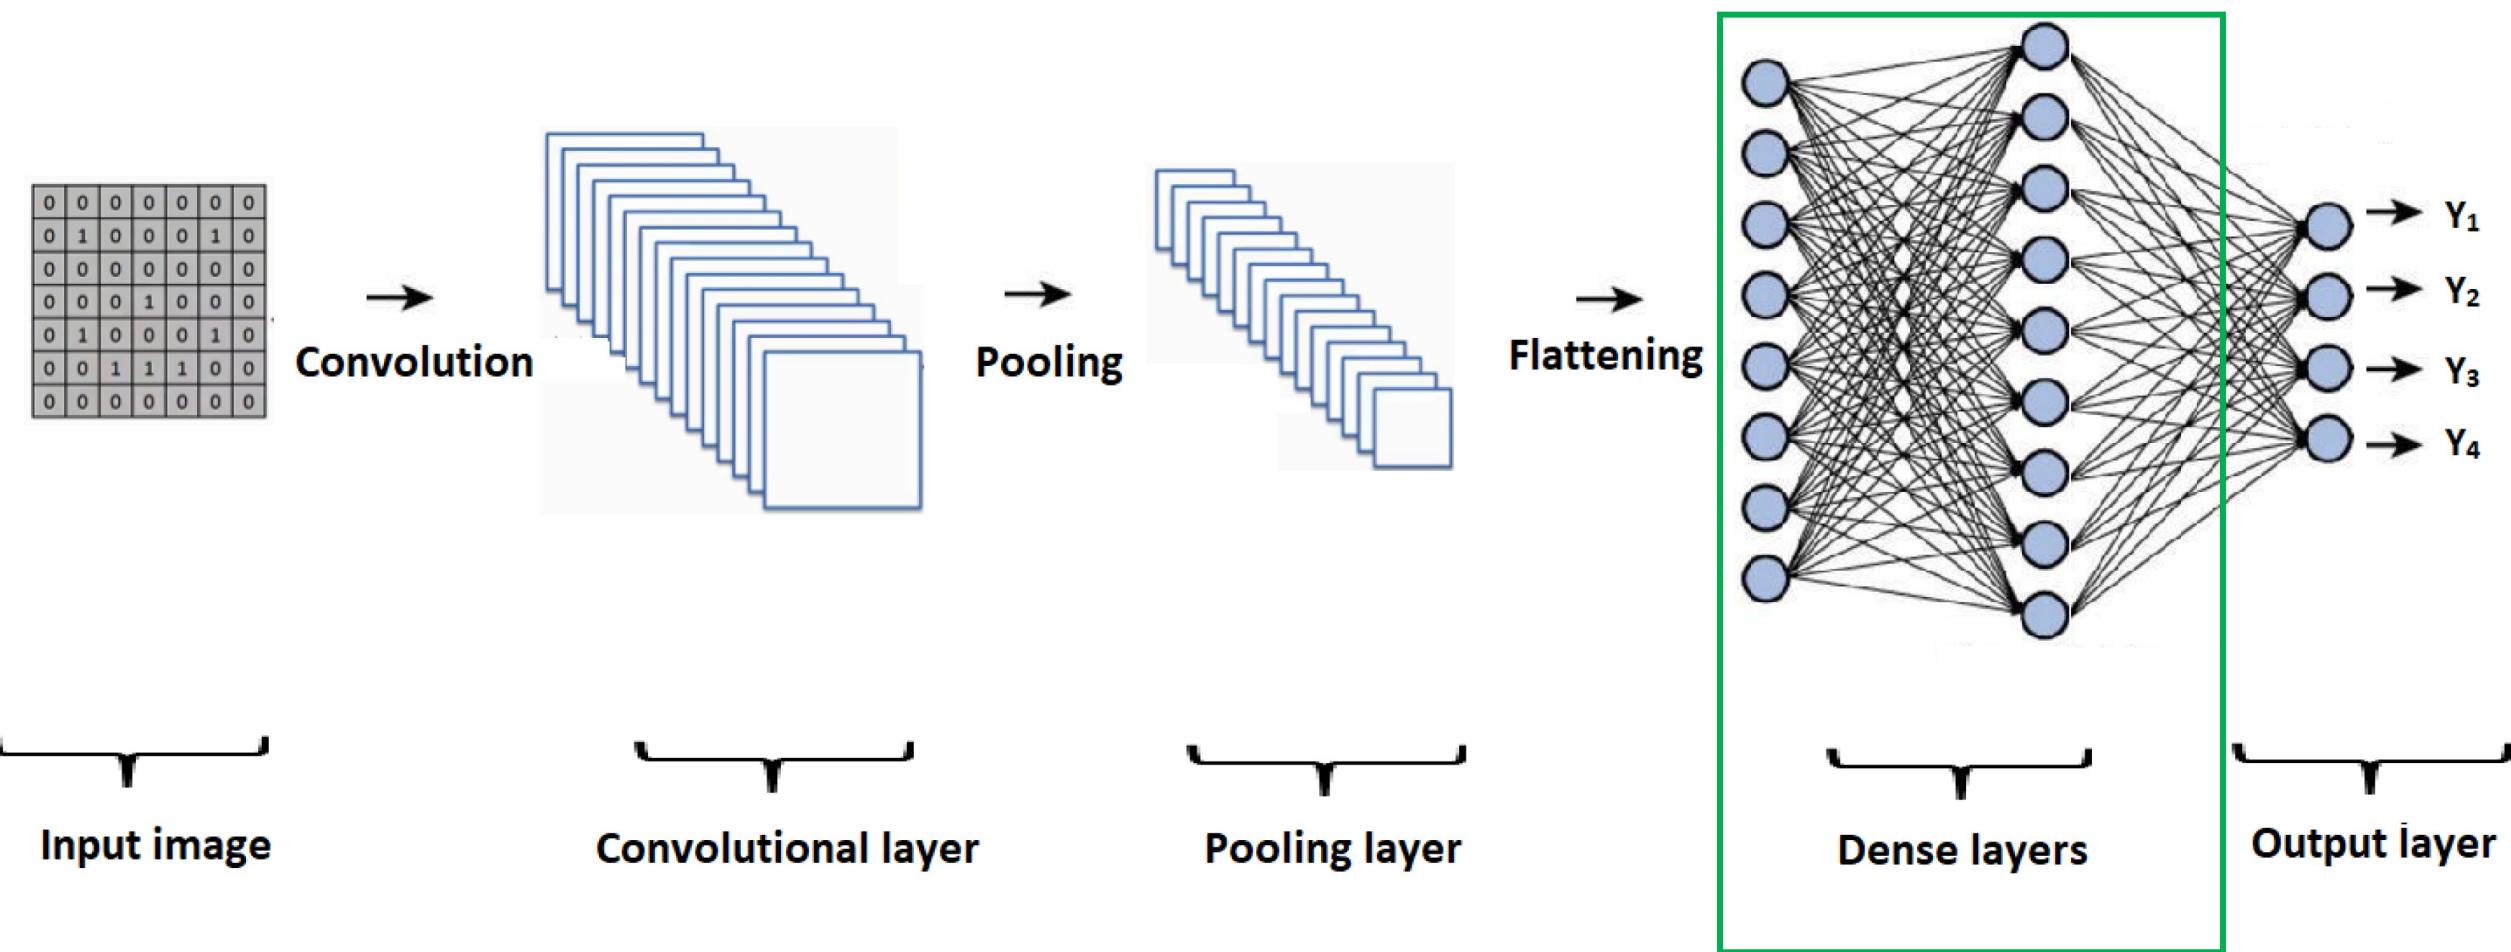
Convolutional Neural Network Architecture



Pooling layer

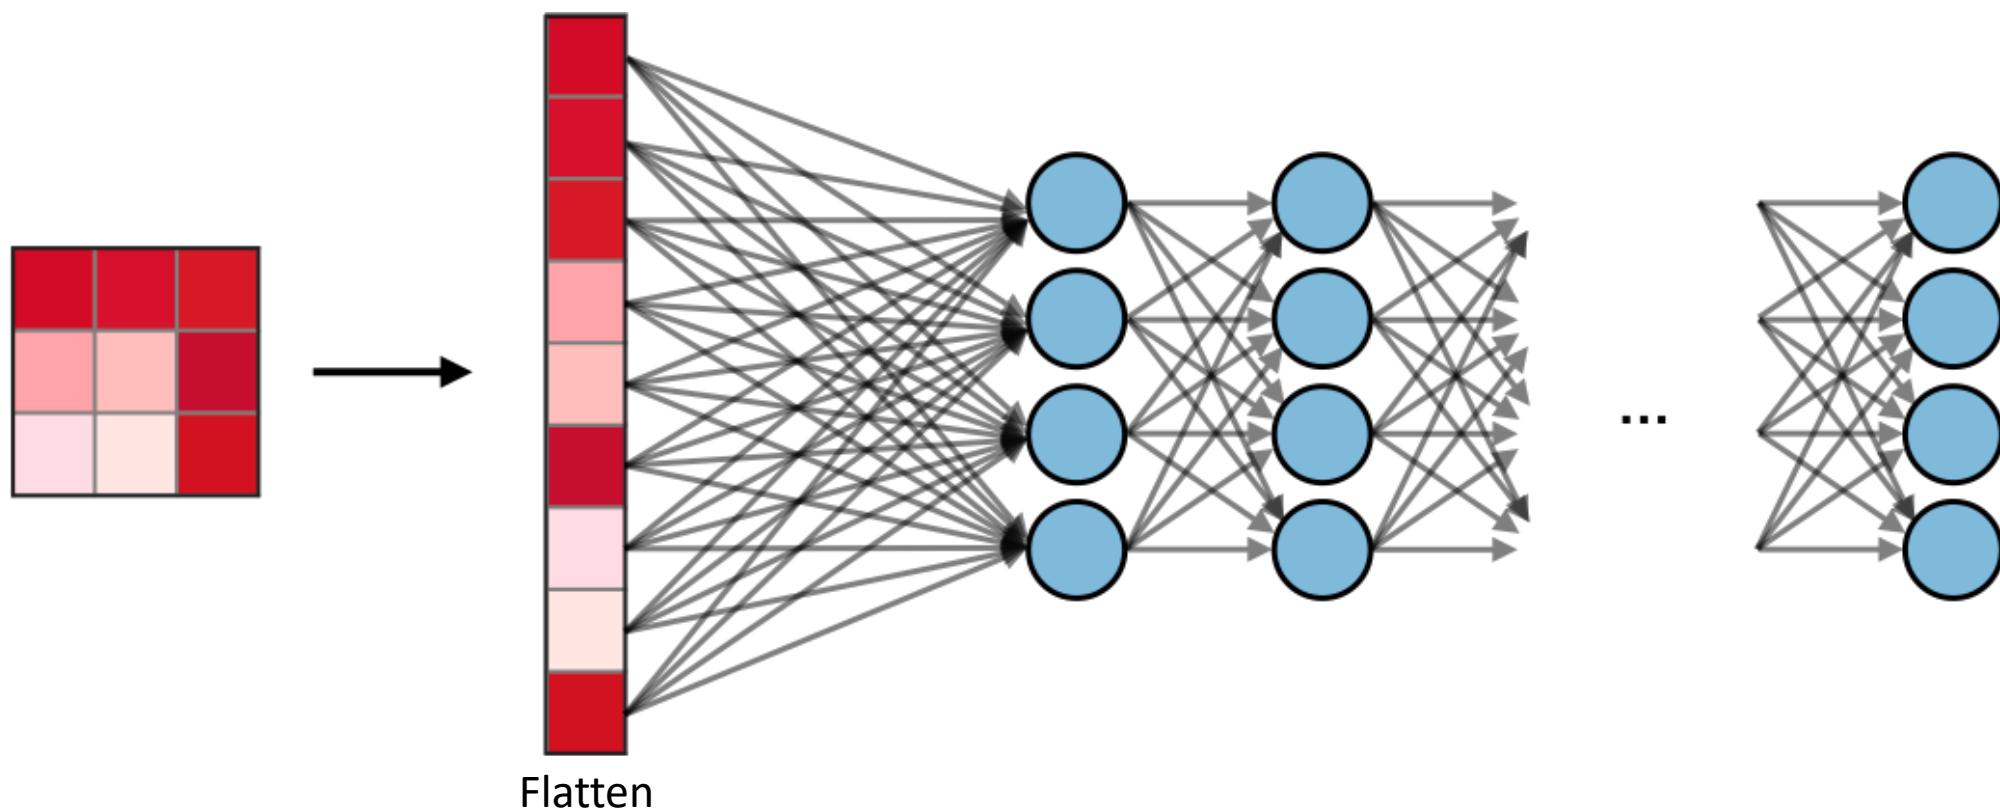


Convolutional Neural Network Architecture

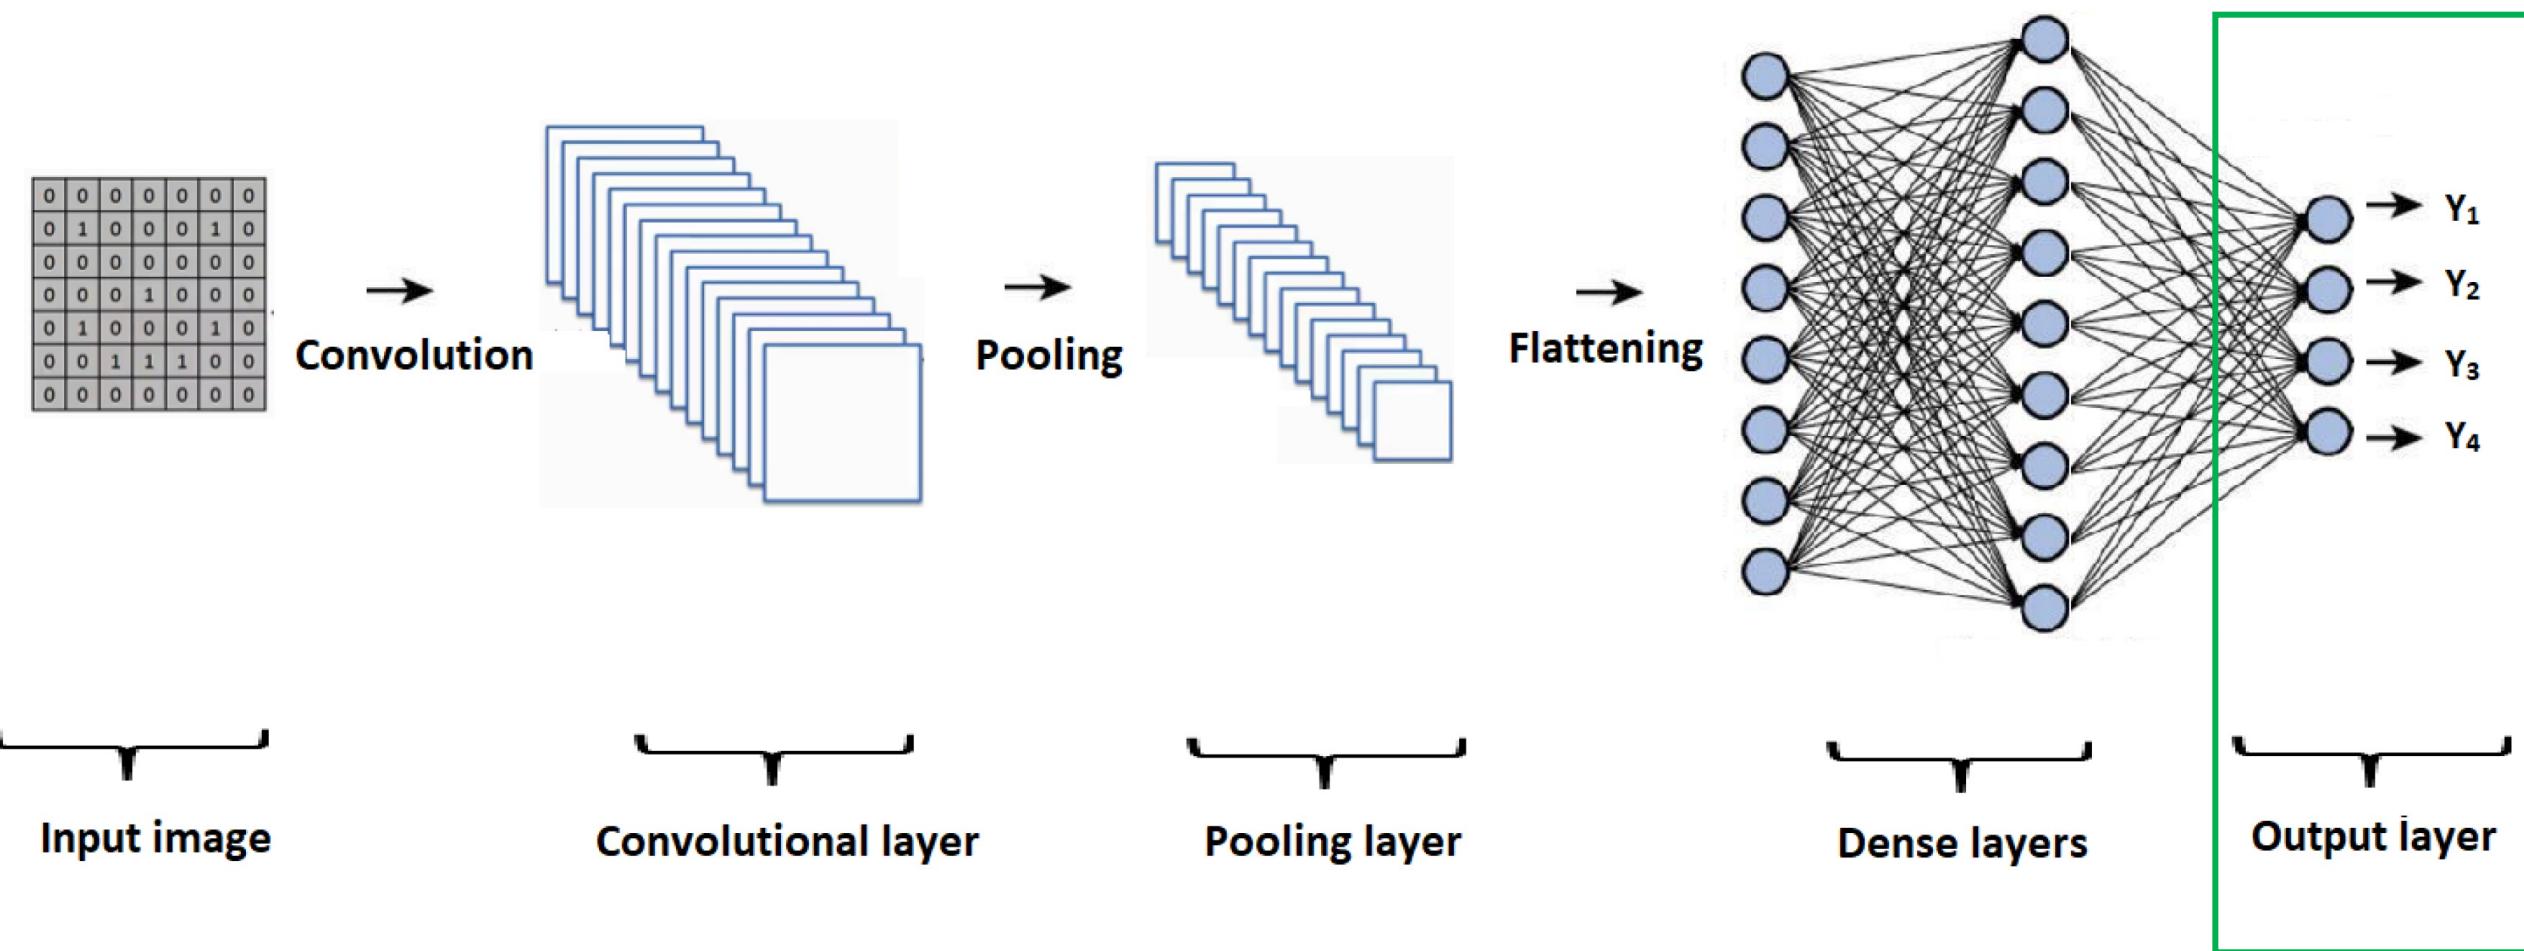


Fully connected layer

- Also known as dense layer
- Each input is connected to all hidden units



Convolutional Neural Network Architecture



Activation functions

- Softmax

$$p = \begin{pmatrix} p_1 \\ \vdots \\ p_n \end{pmatrix}$$

where

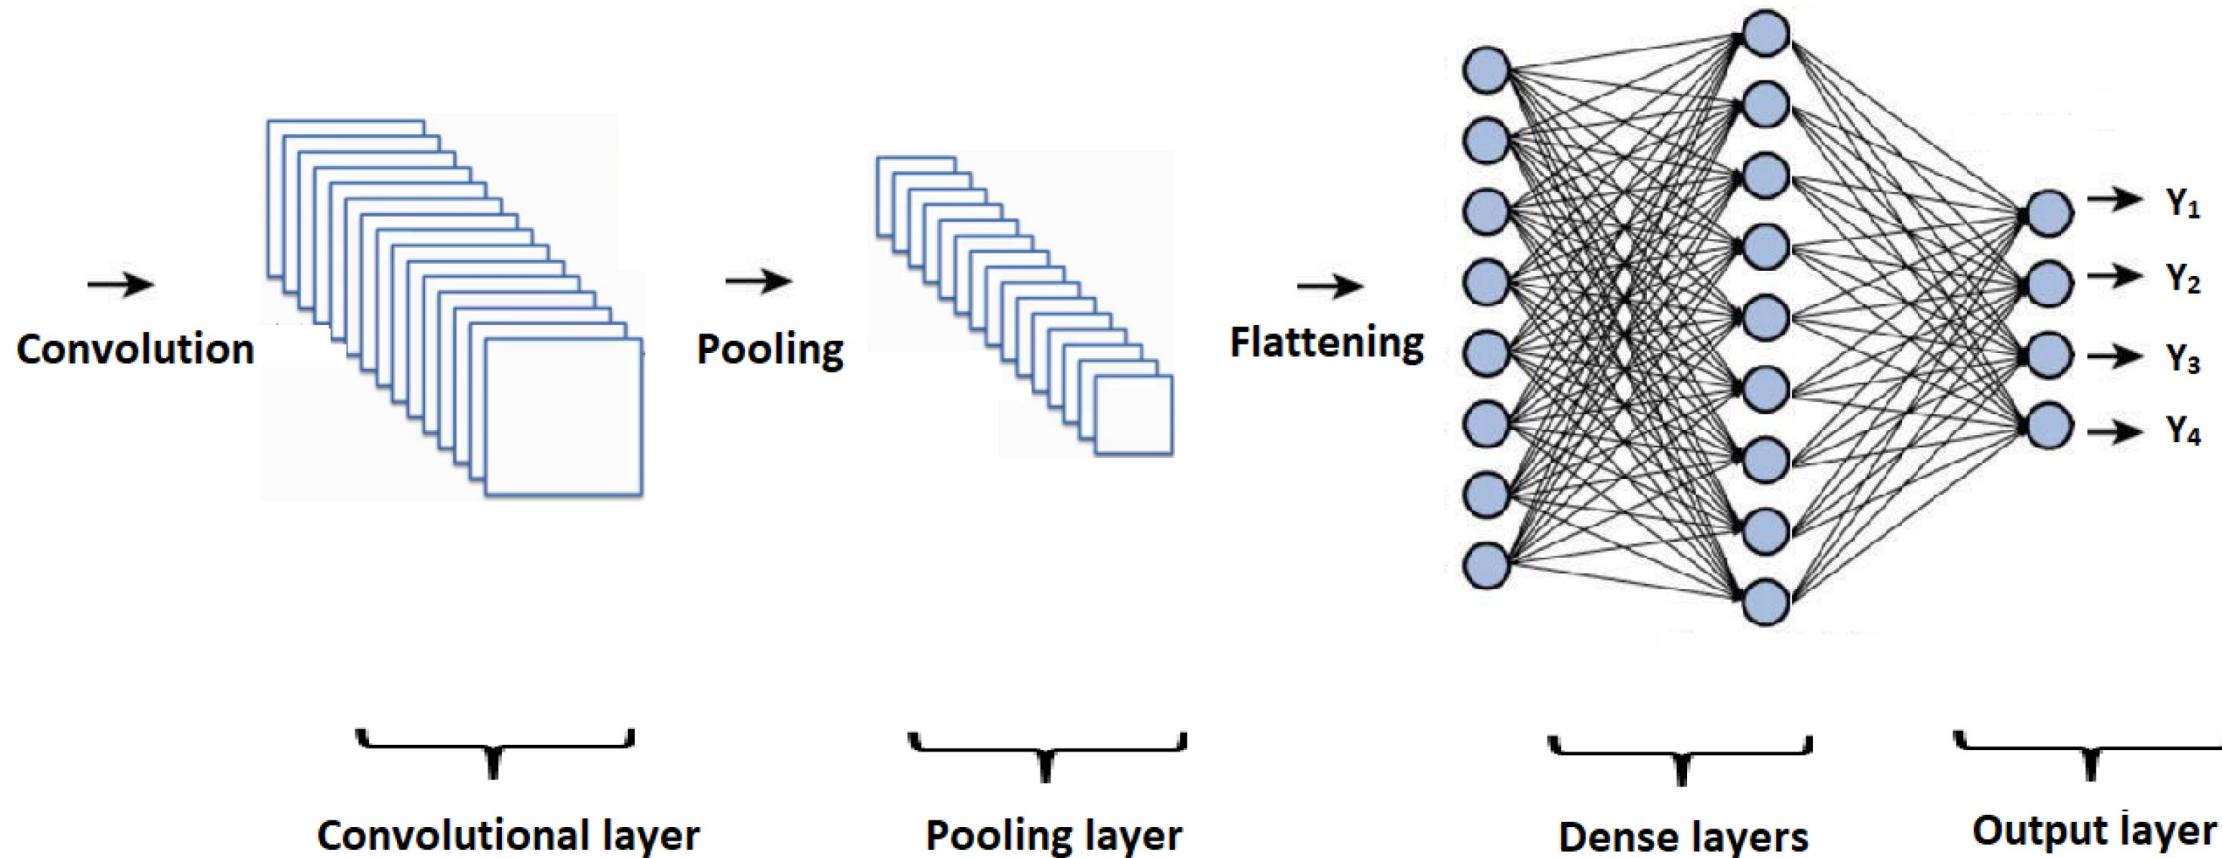
$$p_i = \frac{e^{x_i}}{\sum_{j=1}^n e^{x_j}}$$

- Cross-entropy loss

$$L = -\frac{1}{m} \sum_{i=1}^m y_i \cdot \log(\hat{y}_i)$$

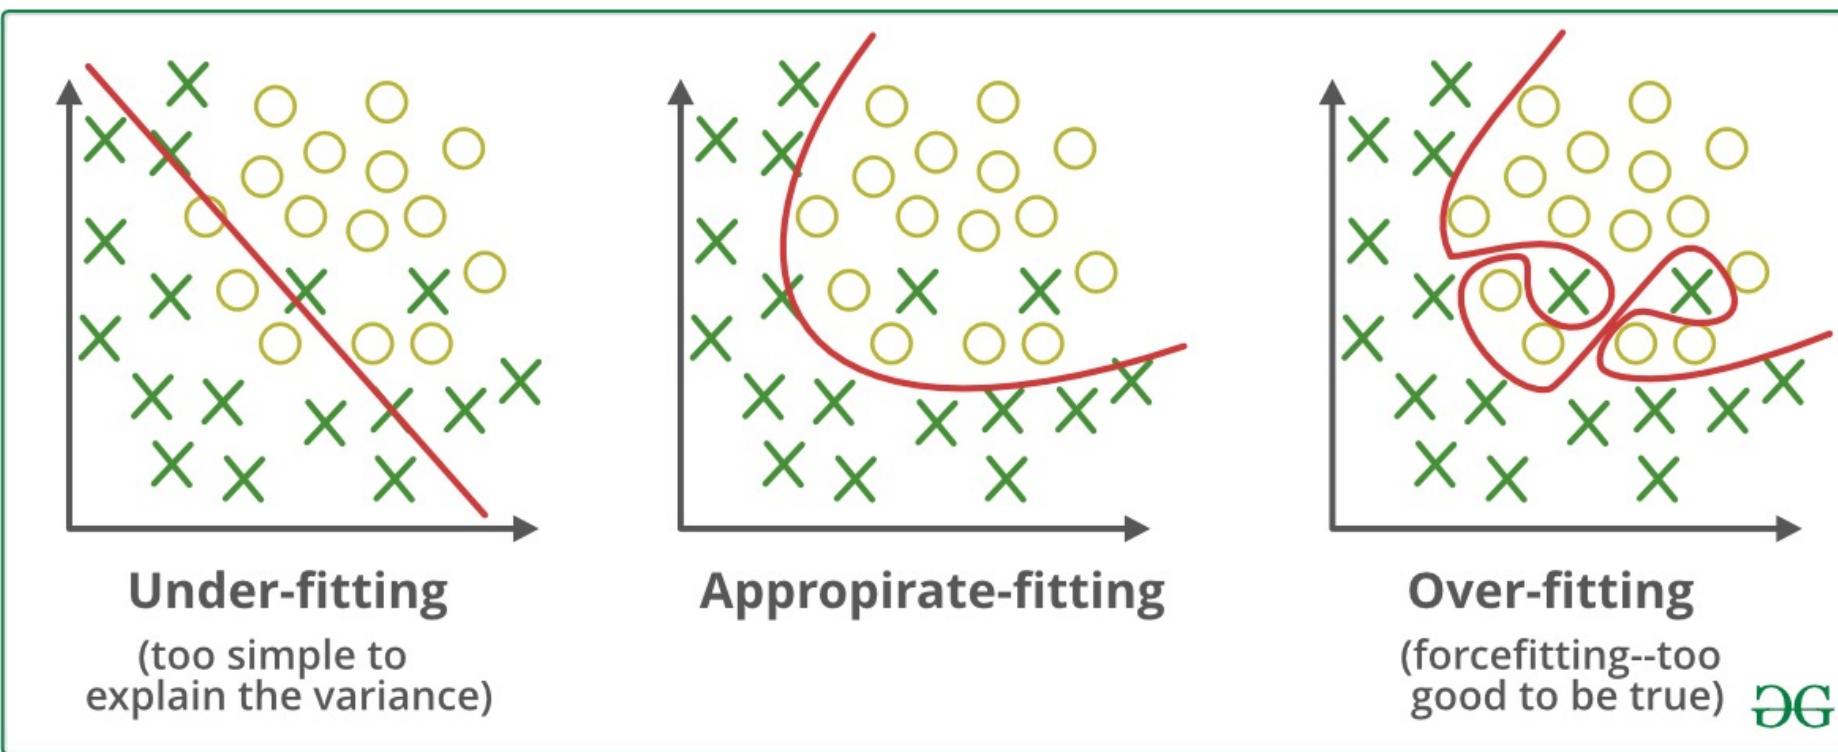
y is a one-hot vector ($[0,0,1,\dots,0]$)

0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	
0	0	0	0	0	0	0	
0	0	0	1	0	0	0	
0	1	0	0	0	1	0	
0	0	1	1	1	0	0	
0	0	0	0	0	0	0	

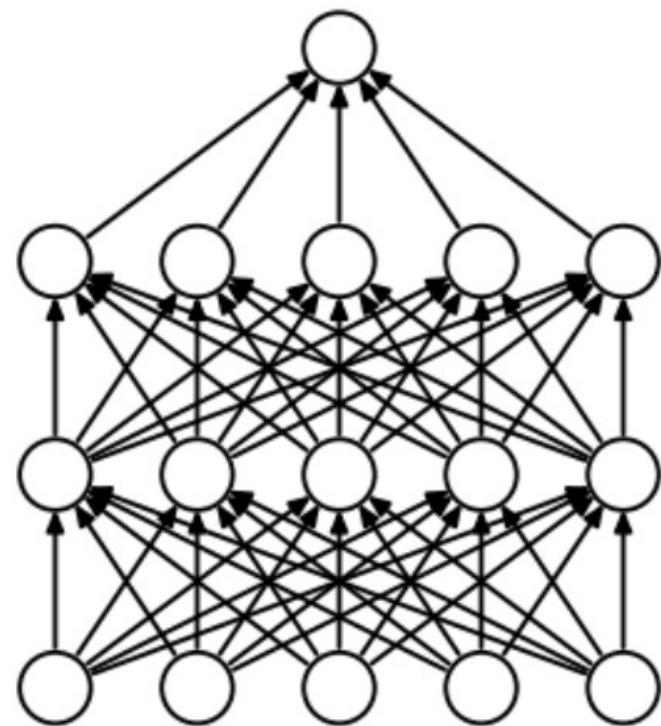


Dropout

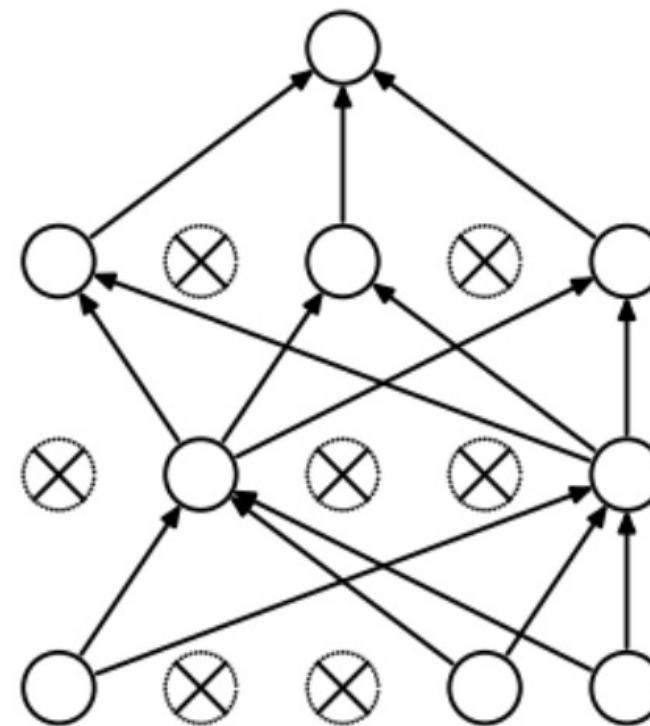
- Overfitting



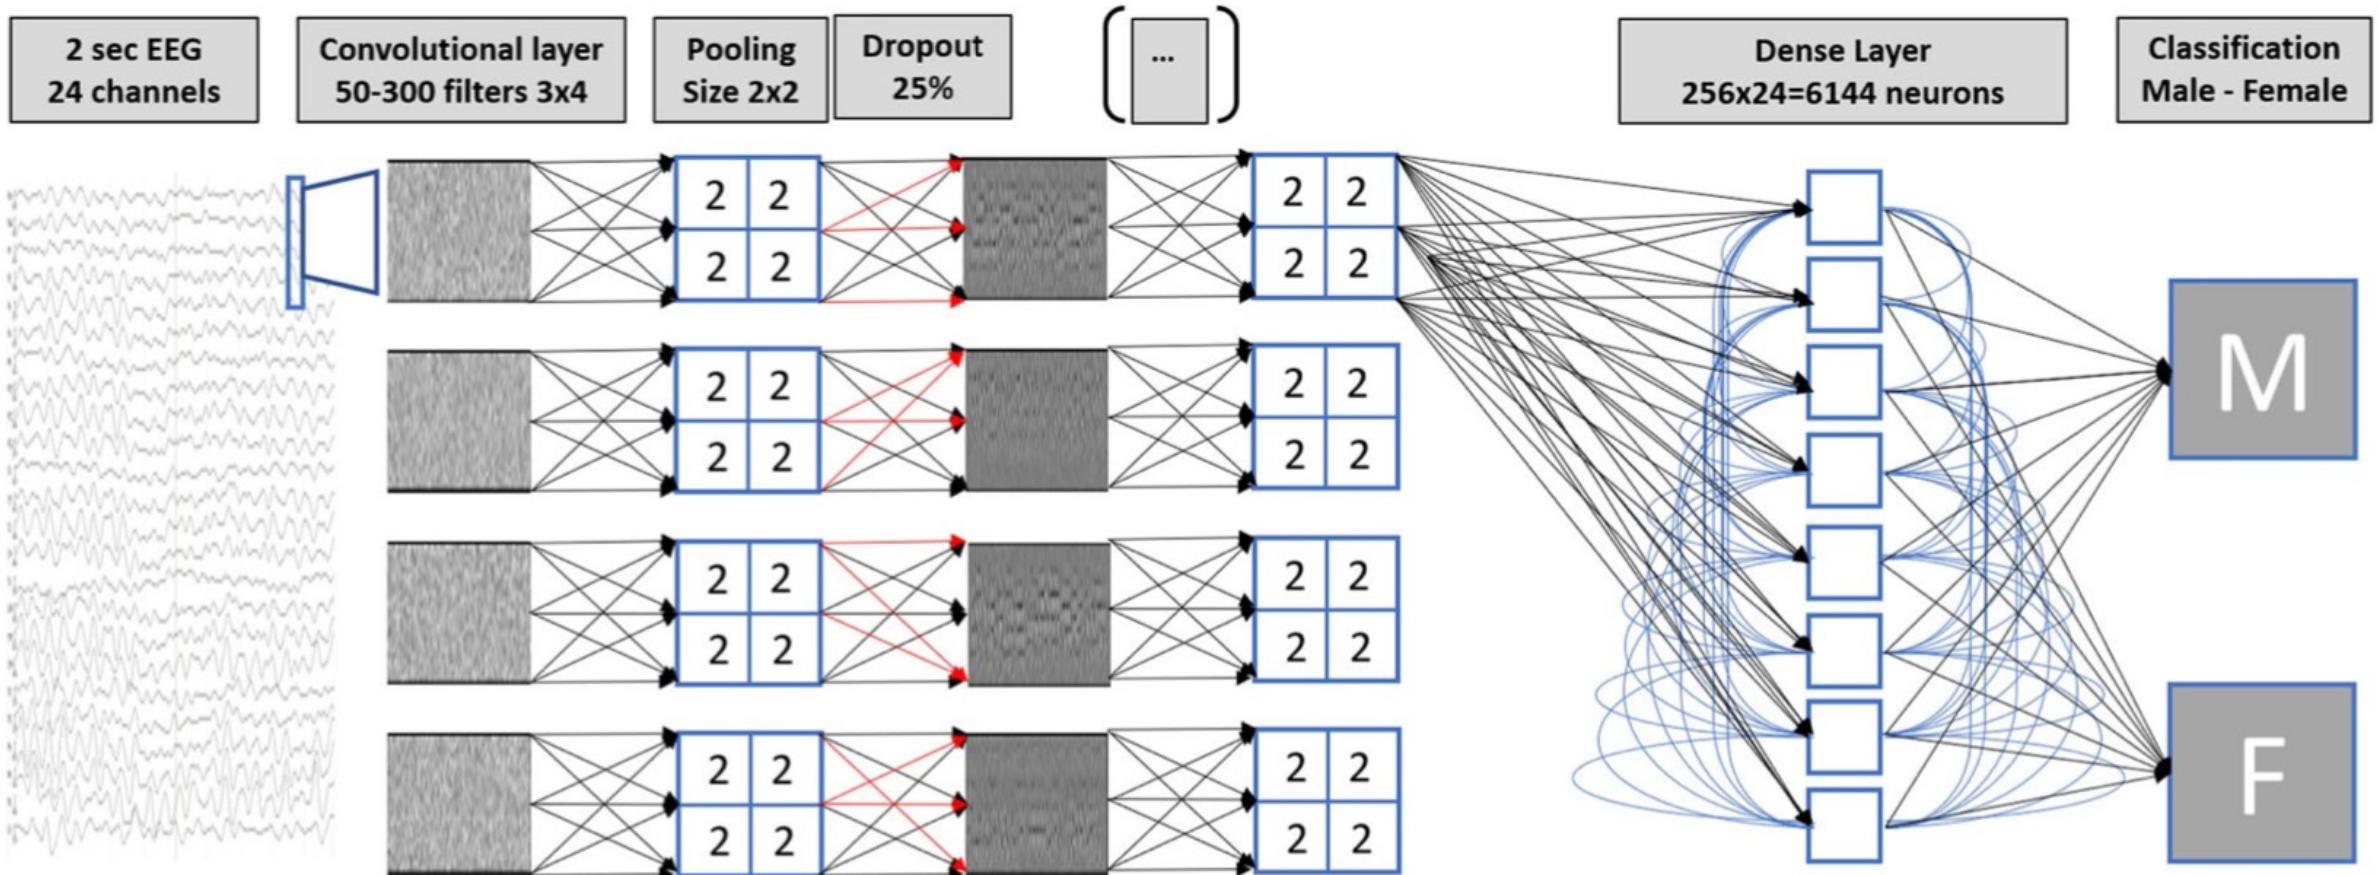
Dropout



(a) Standard Neural Net



(b) After applying dropout.



The model

Layer (type)	Input Shape	Param #	Tr. Param #
=====			
Conv2d-1	[1, 1, 24, 256]	1,000	1,000
ReLU-2	[1, 100, 22, 254]	0	0
MaxPool2d-3	[1, 100, 22, 254]	0	0
Dropout-4	[1, 100, 11, 127]	0	0
Conv2d-5	[1, 100, 11, 127]	90,100	90,100
ReLU-6	[1, 100, 9, 125]	0	0
MaxPool2d-7	[1, 100, 9, 125]	0	0
Dropout-8	[1, 100, 4, 62]	0	0
Conv2d-9	[1, 100, 4, 62]	180,300	180,300
ReLU-10	[1, 300, 3, 60]	0	0
MaxPool2d-11	[1, 300, 3, 60]	0	0
Dropout-12	[1, 300, 1, 30]	0	0
Conv2d-13	[1, 300, 1, 30]	630,300	630,300
ReLU-14	[1, 300, 1, 24]	0	0
MaxPool2d-15	[1, 300, 1, 24]	0	0
Dropout-16	[1, 300, 1, 23]	0	0
Conv2d-17	[1, 300, 1, 23]	90,100	90,100
Conv2d-18	[1, 100, 1, 21]	30,100	30,100
Flatten-19	[1, 100, 1, 19]	0	0
Linear-20	[1, 1900]	11,679,744	11,679,744
Linear-21	[1, 6144]	12,290	12,290
=====			
Total params: 12,713,934			
Trainable params: 12,713,934			
Non-trainable params: 0			
=====			